

Microservices

...

Calin-loan Nicoara

Senior Software Developer and Trainer



Let's get to know each other



Outline

- Microservice architecture
- Intercommunication
- Security
- Testing strategies
- Deployment
- Monitoring

What are Microservices?

A microservice application is a collection of independent services, each doing one thing well, that work together to perform more complex operations.

Before Microservices

Monolithic

Tight coupling



All elements must agree on each change

Traditional SOA

Looser coupling



Elements must be coordinated each other

Microservice

Decoupled



Elements can be changed or added without prior coordination with others

Short history of microservices

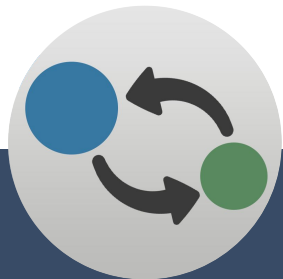


Microservice benefits



Autonomy

Autonomous team
with isolated
implementation



Speed of change

Build and deploy
small independent
services



Scale

Functional and
Team Scaling



Tech Diversity

Allows try-out of
new technologies



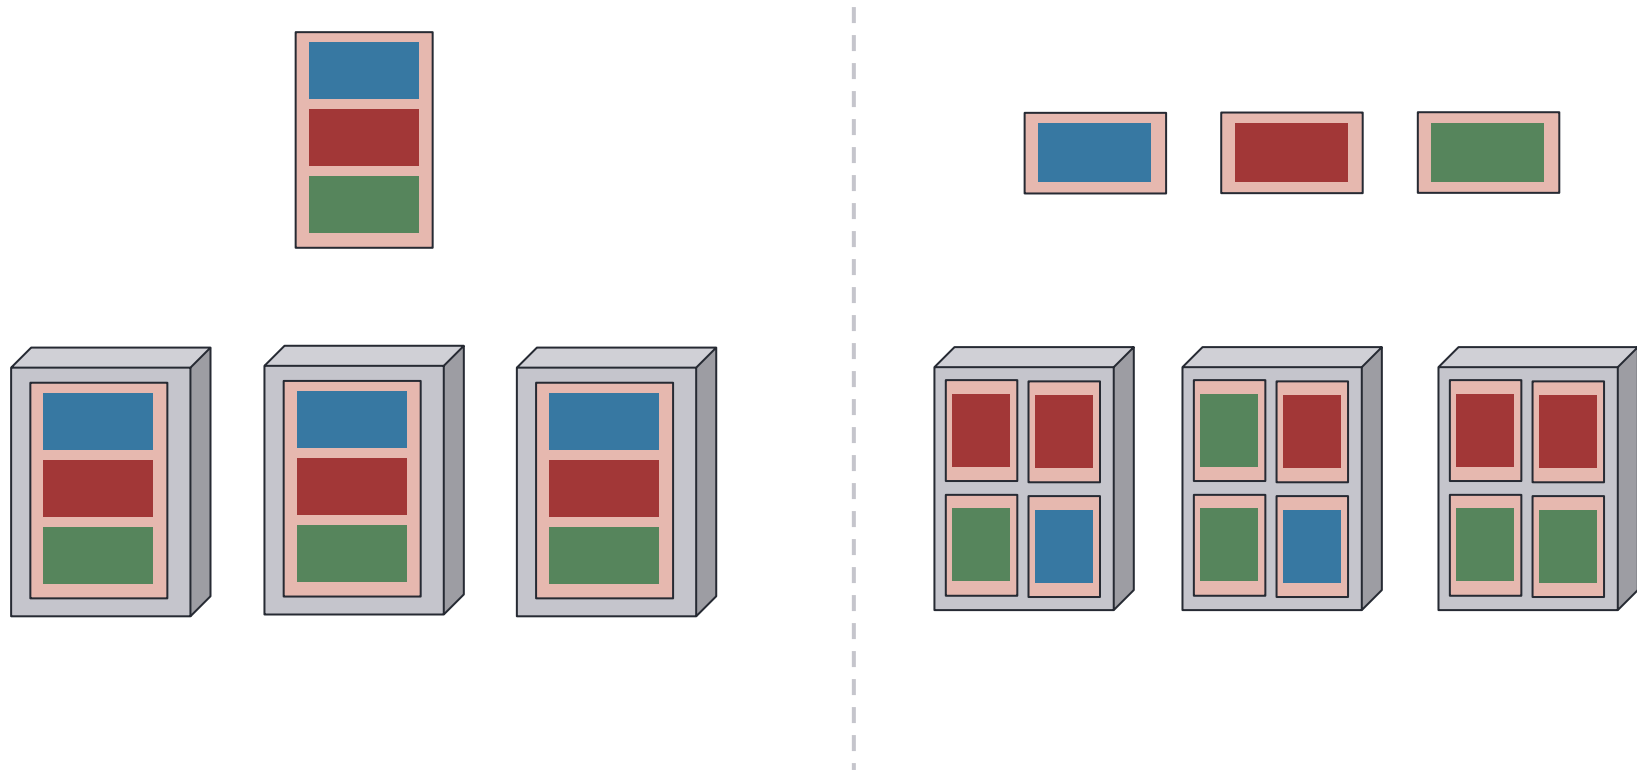
Intelligible

Lower cognitive
load

Microservices

the main characteristics

Scalability & High-availability



Organized around business capabilities



UI



Server



DBA



Shipping

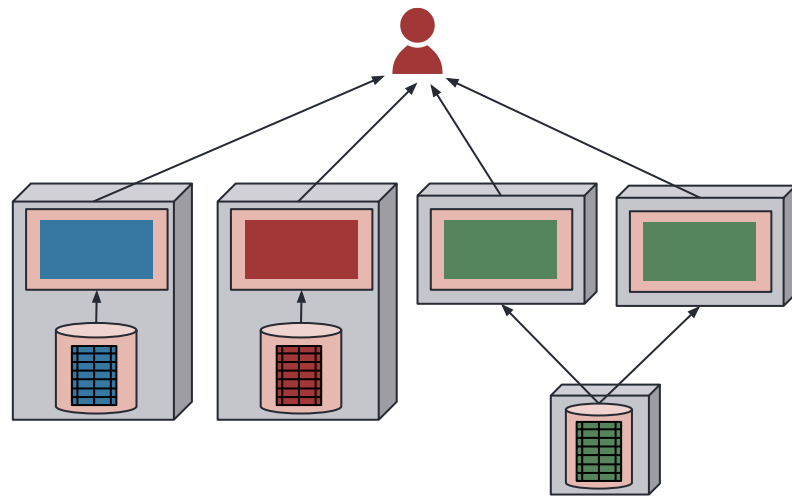
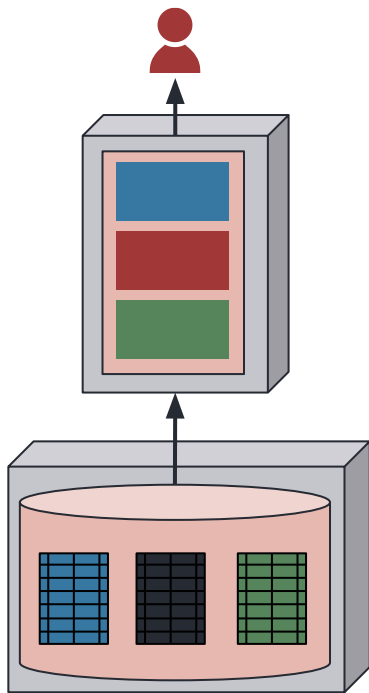


Orders



Catalog

Autonomous (from design to deployment)



Disadvantages of distributed computing

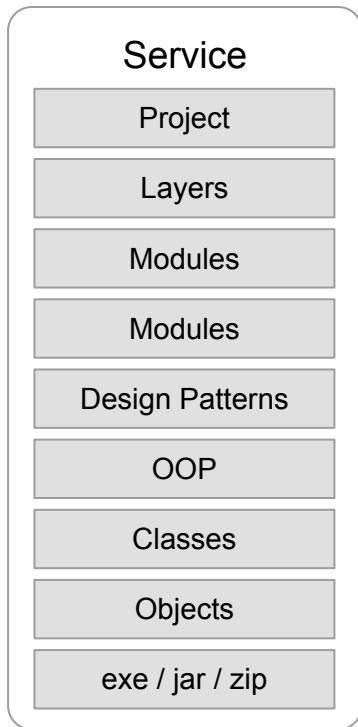
- Latency
- Partial failure
- Infrastructure

Designing Microservices

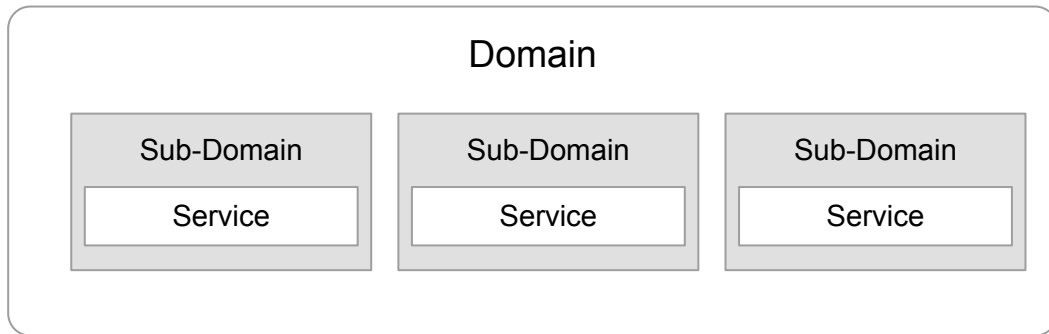
- **Size**
Single Responsibility Principle
- **Scope**
Agile / Independent development
and deployment
- **Capabilities**
Bounded Context in
Domain-Driven-Design

Domain Driven Design

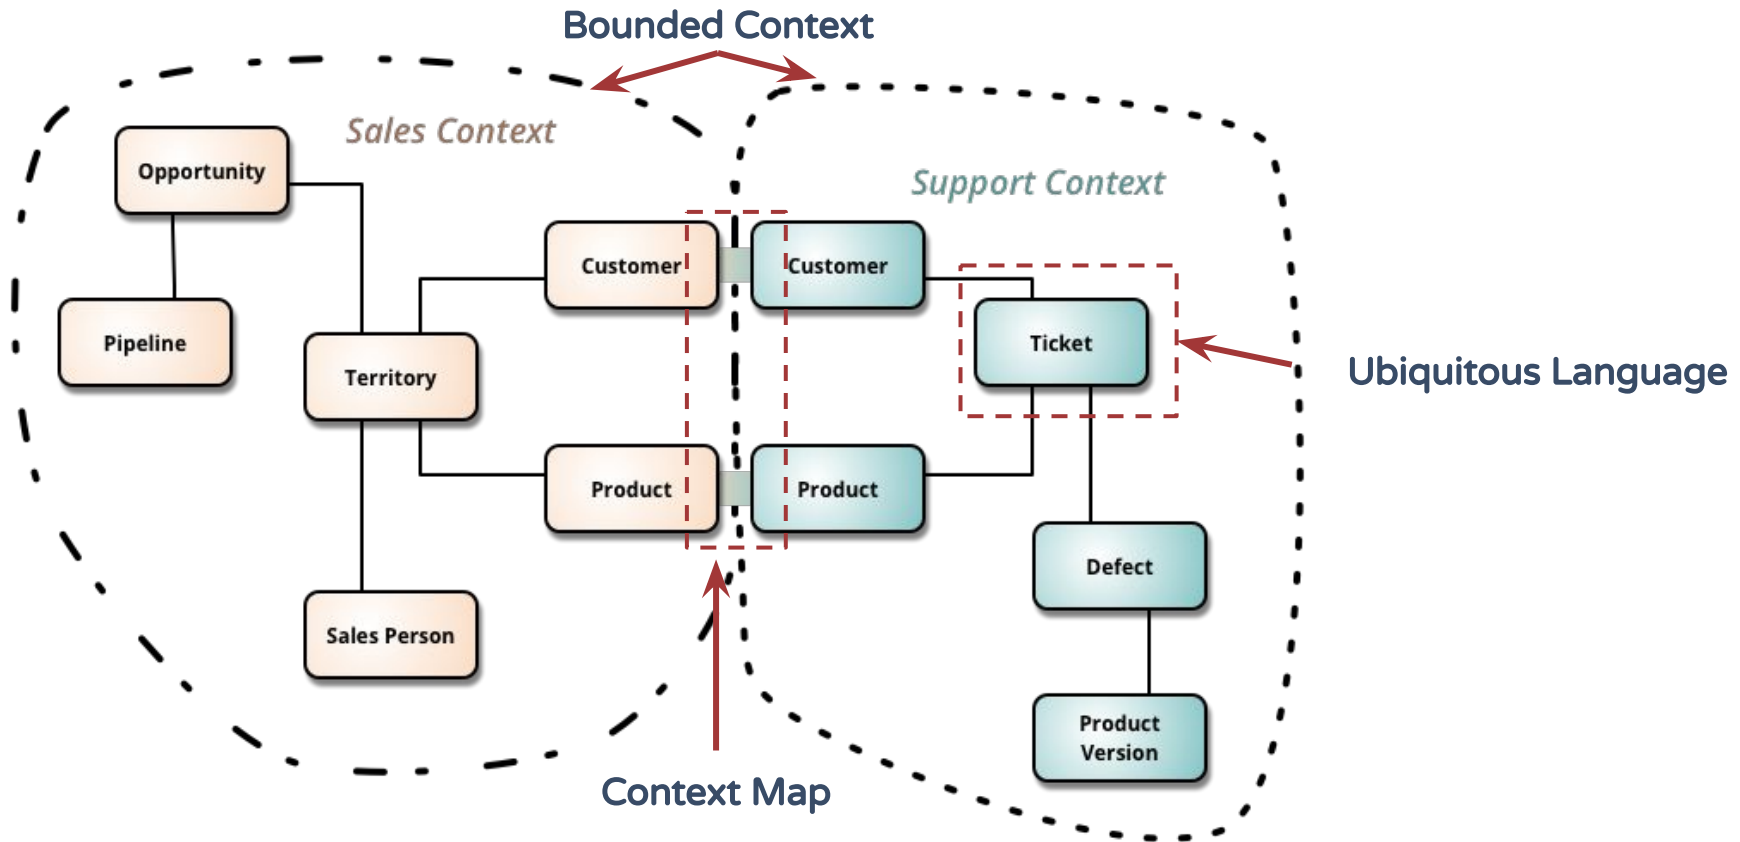
Tactical Design



Strategic Design



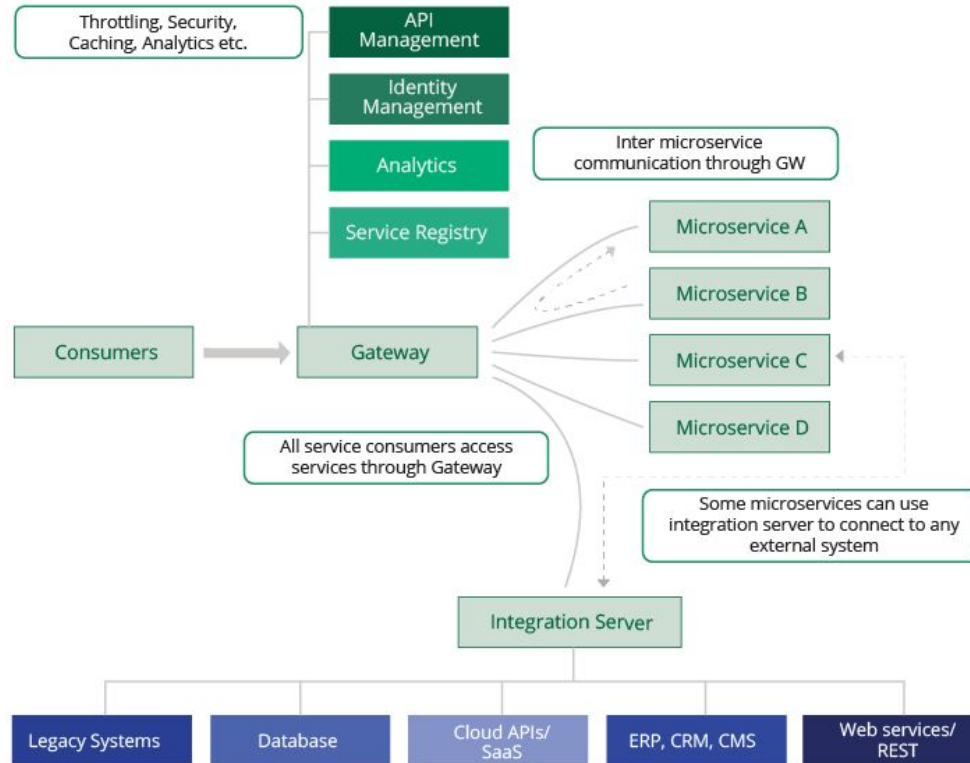
DDD - Strategic Design



Microservices in Practice

eCommerce

Microservices Architecture



eCommerce Microservices

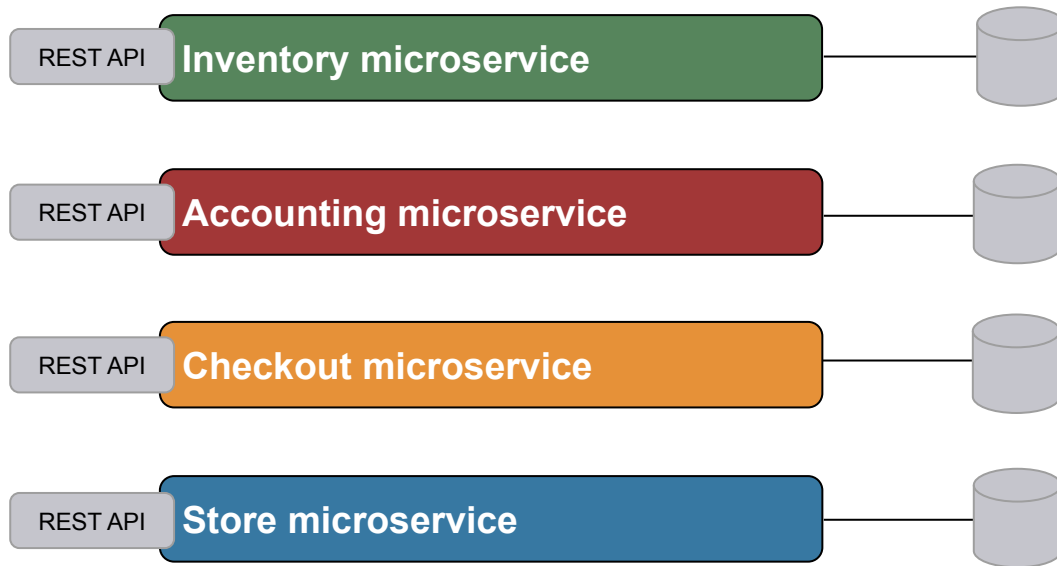
Inventory microservice

Accounting microservice

Checkout microservice

Store microservice

eCommerce Microservices



REST

HTTP verbs

- Get
- Post
- Put
- Delete
- Patch

Business Capabilities

Inventory Context

Pricing

Stock

Accounting Context

Customer

Address

Store Context

Product

Category

Checkout Context

Cart

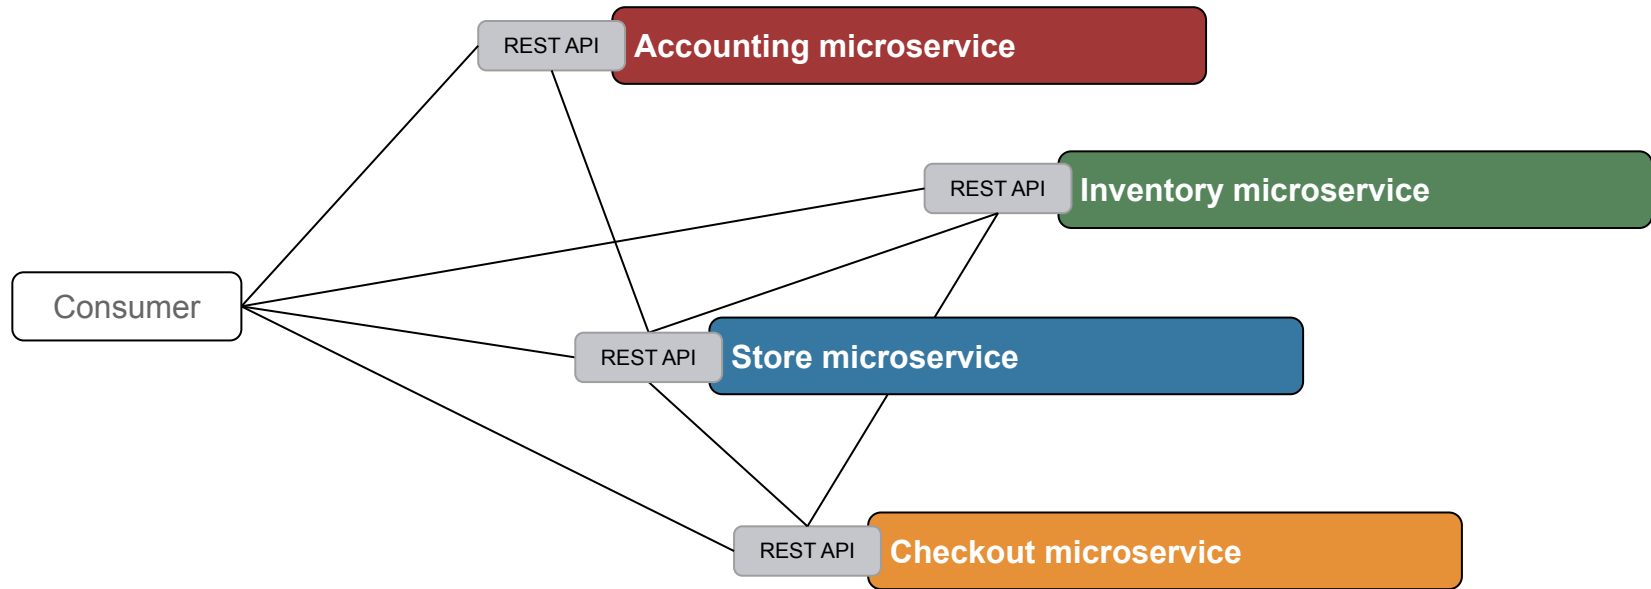
Order

Communication

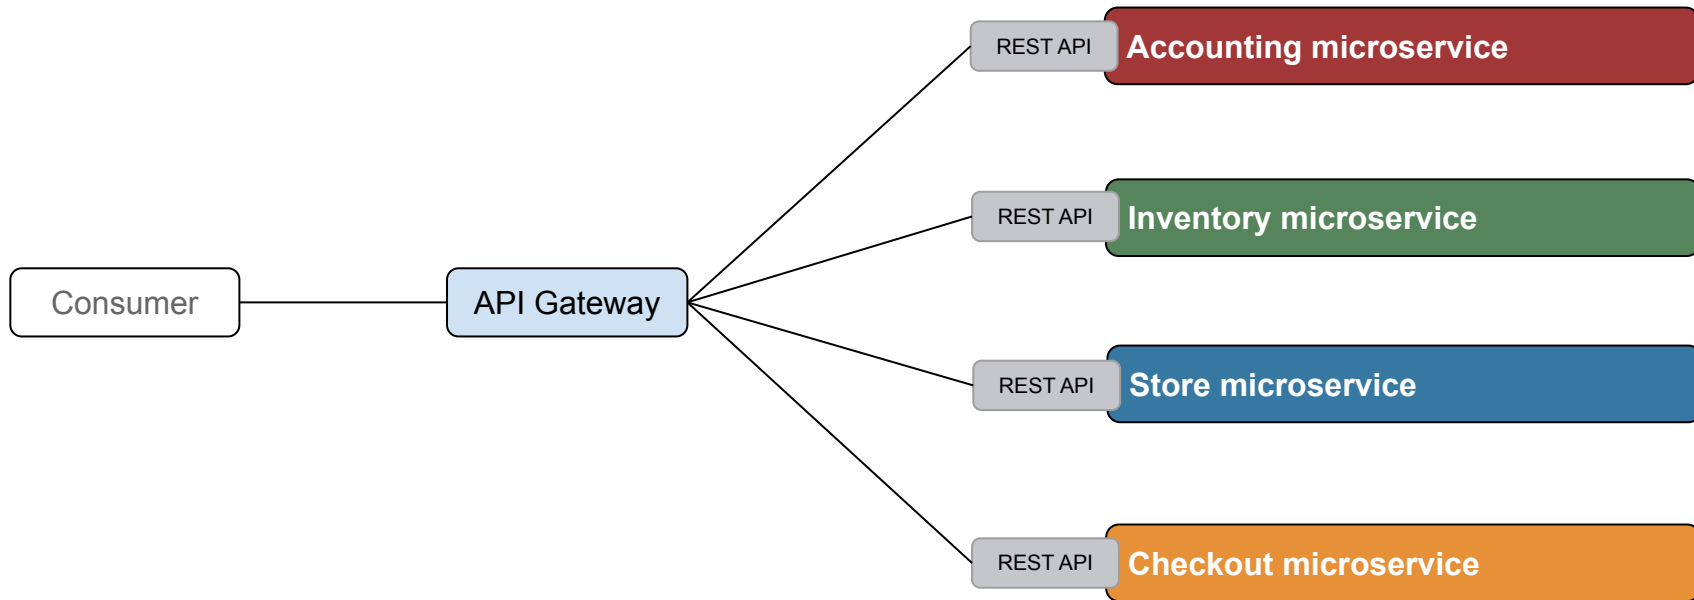
Interaction Styles

| | One-to-One | One-to-Many |
|--------------|---|--|
| Synchronous | Request / Response | -- |
| Asynchronous | Notification ----- Request / Async response | Publish / Subscribe ----- Publish / Async response |

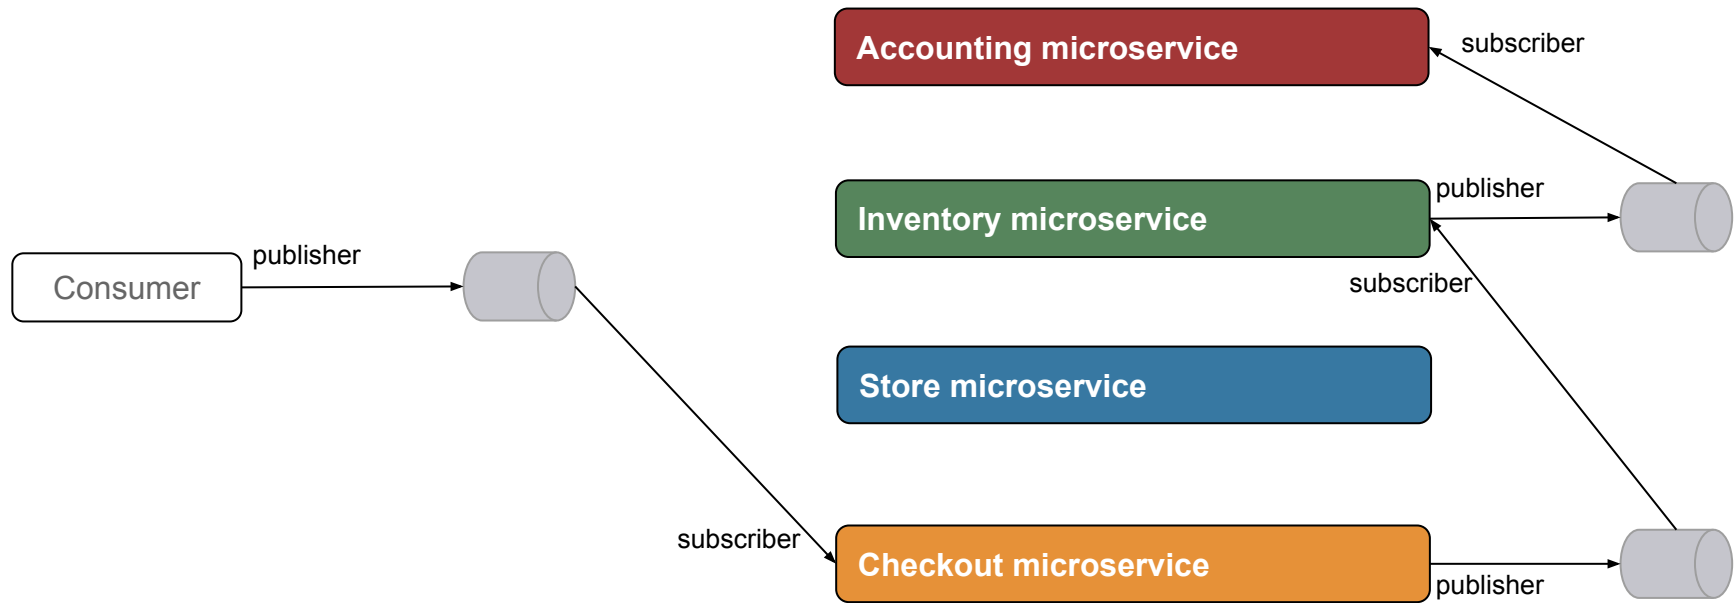
Point-to-point style



Api-Gateway style

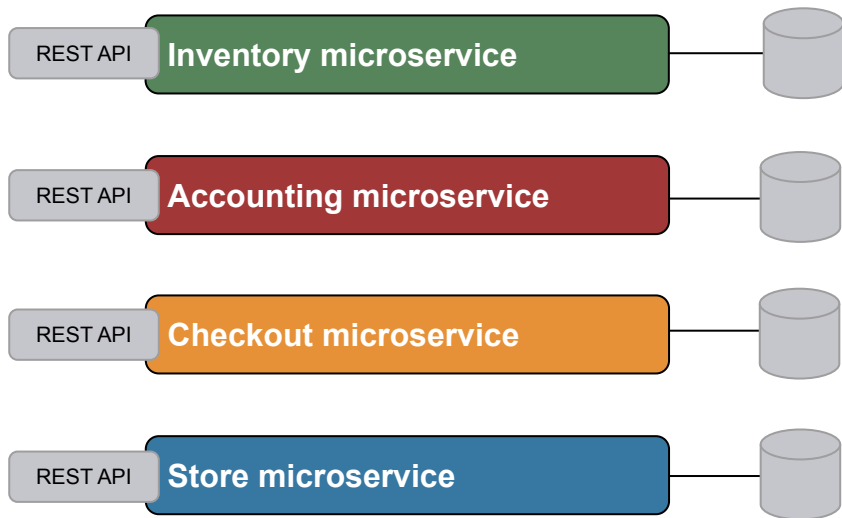


Message-Broker style



Coding time

Coding time



Inventory Context

Pricing

Stock

Accounting Context

Customer

Address

Store Context

Product

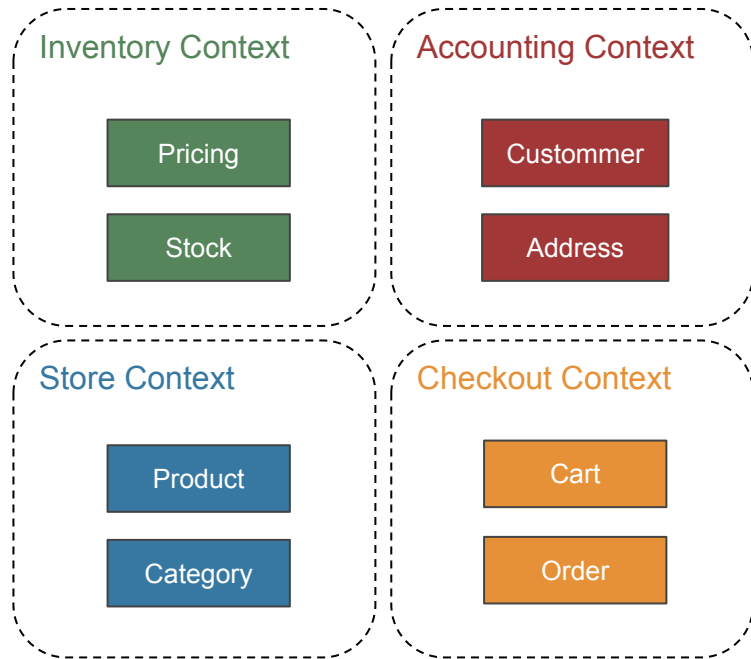
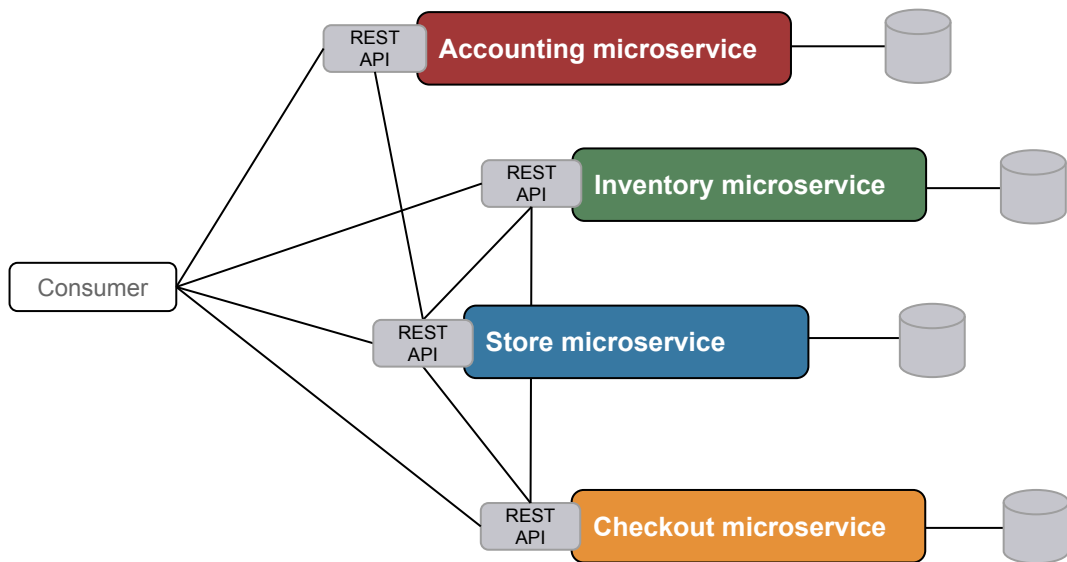
Category

Checkout Context

Cart

Order

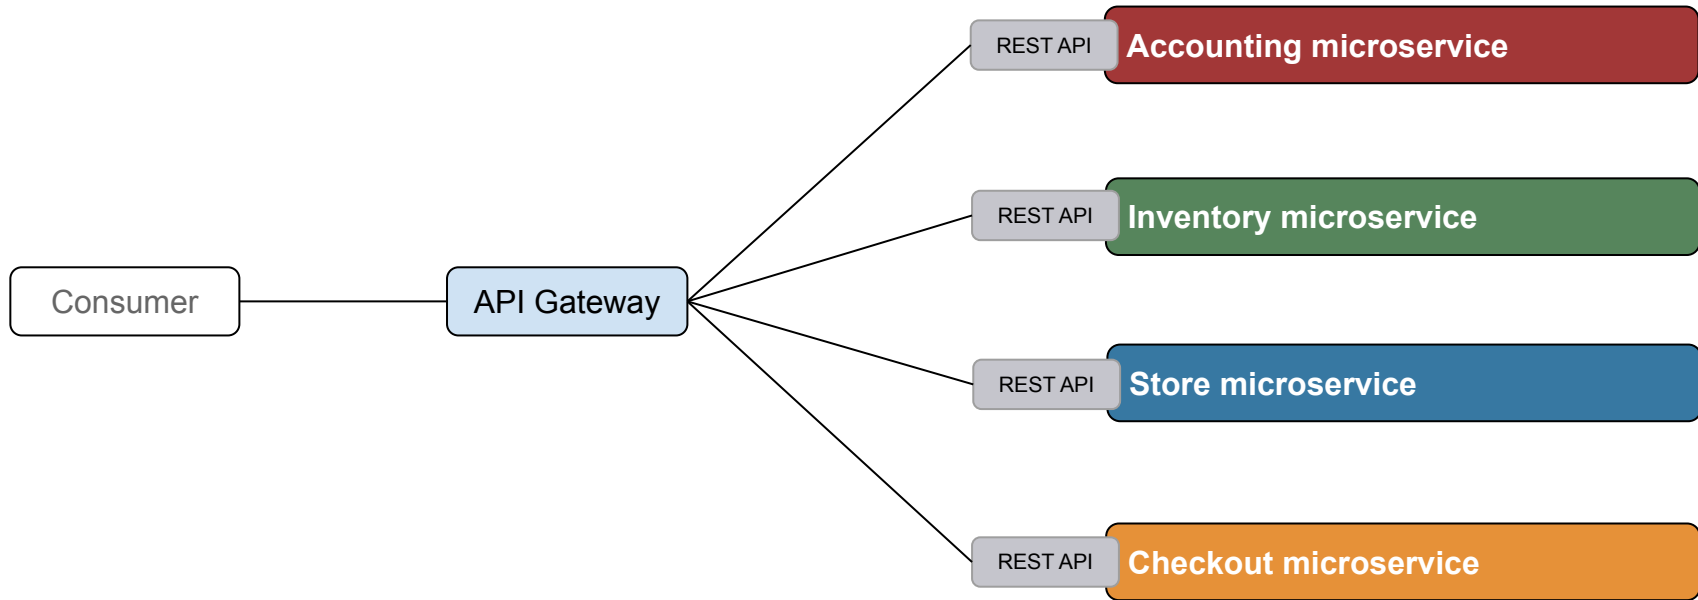
Coding time



Spring Cloud Gateway

- Front door
- Edge service
- Routing and Filtering
- Monitoring
- Security

Coding time



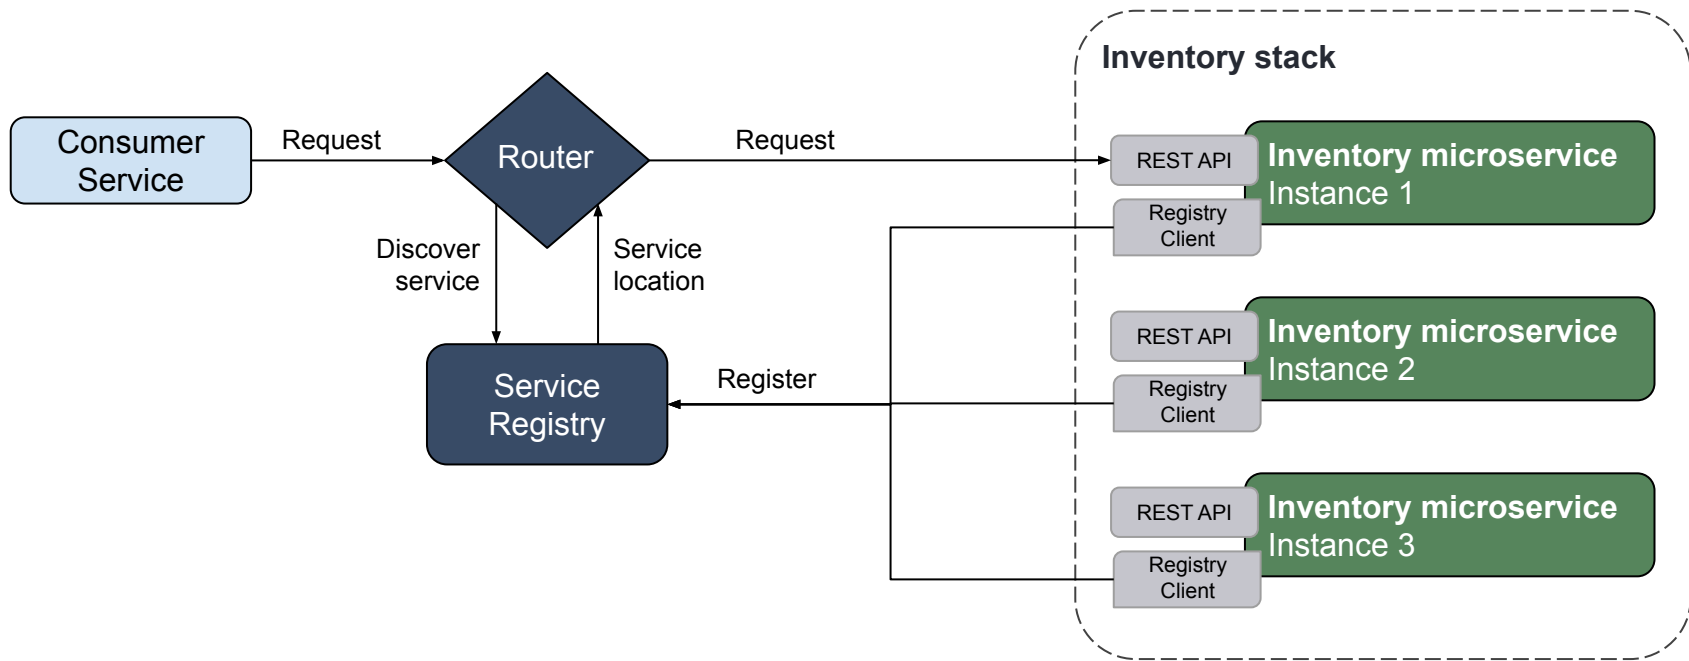
Service Discovery

Server-side and Client-side

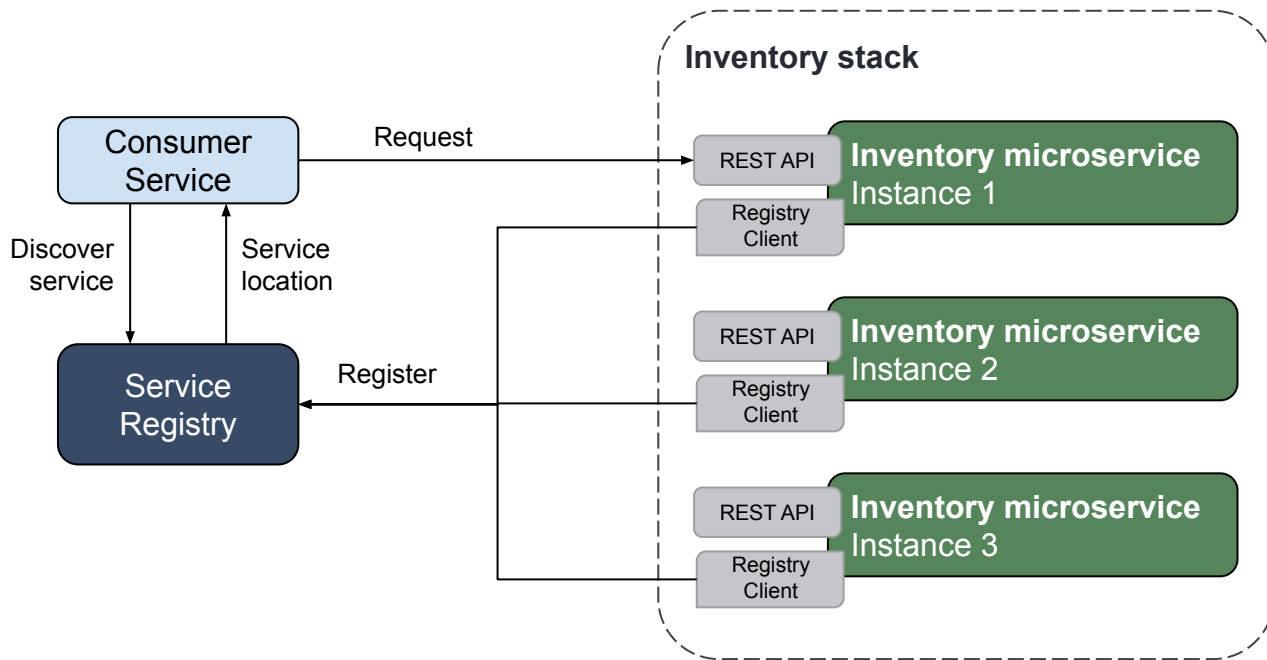
Why use Service Discovery

- Dynamically network location
- Autoscaling
- Failures
- Upgrades

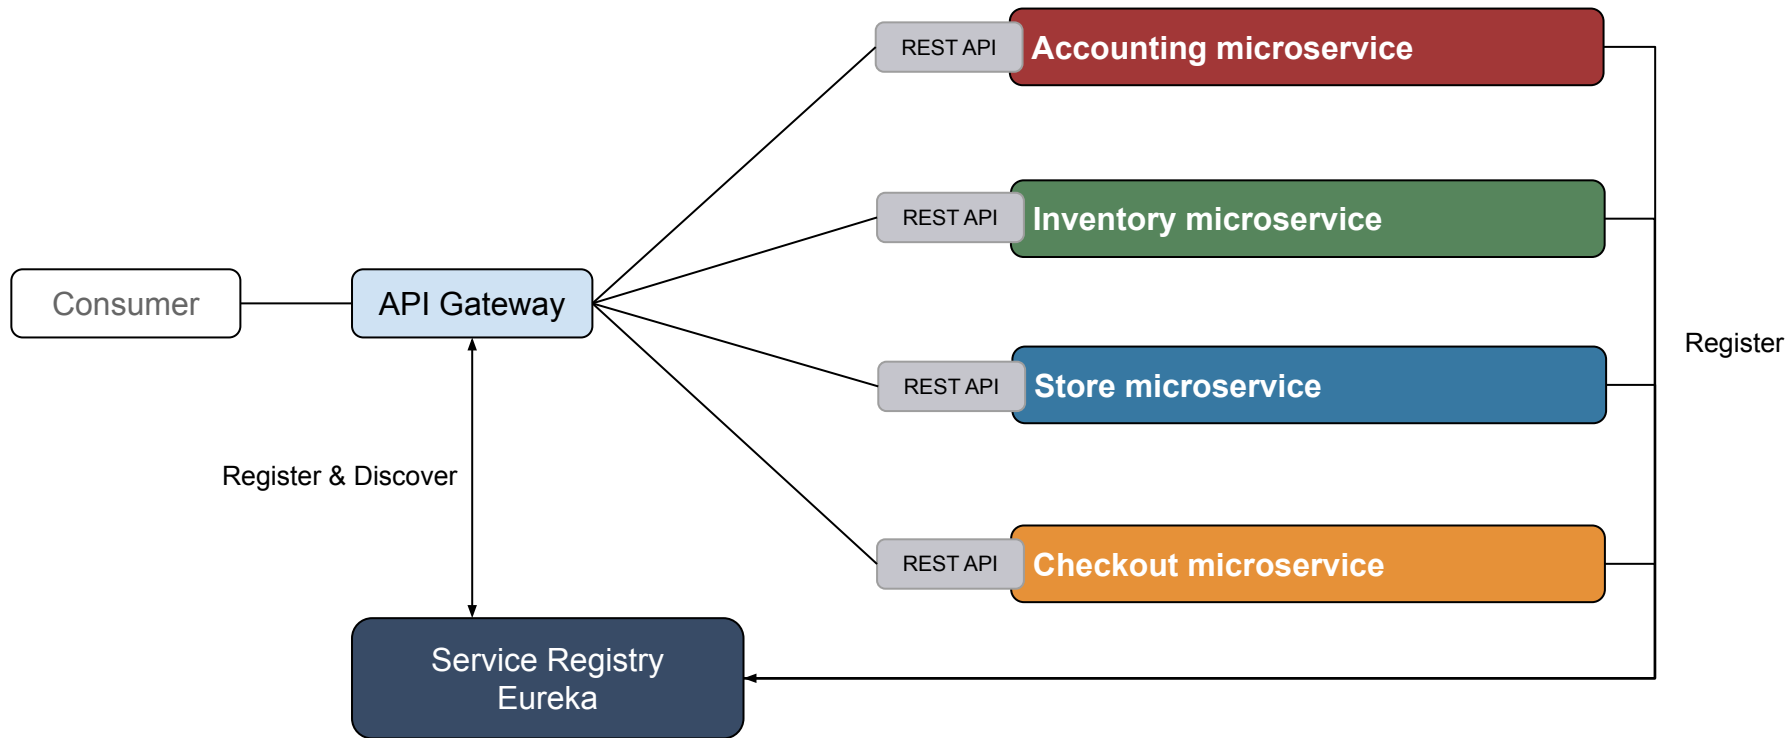
Server-Side Discovery



Client-Side Discovery



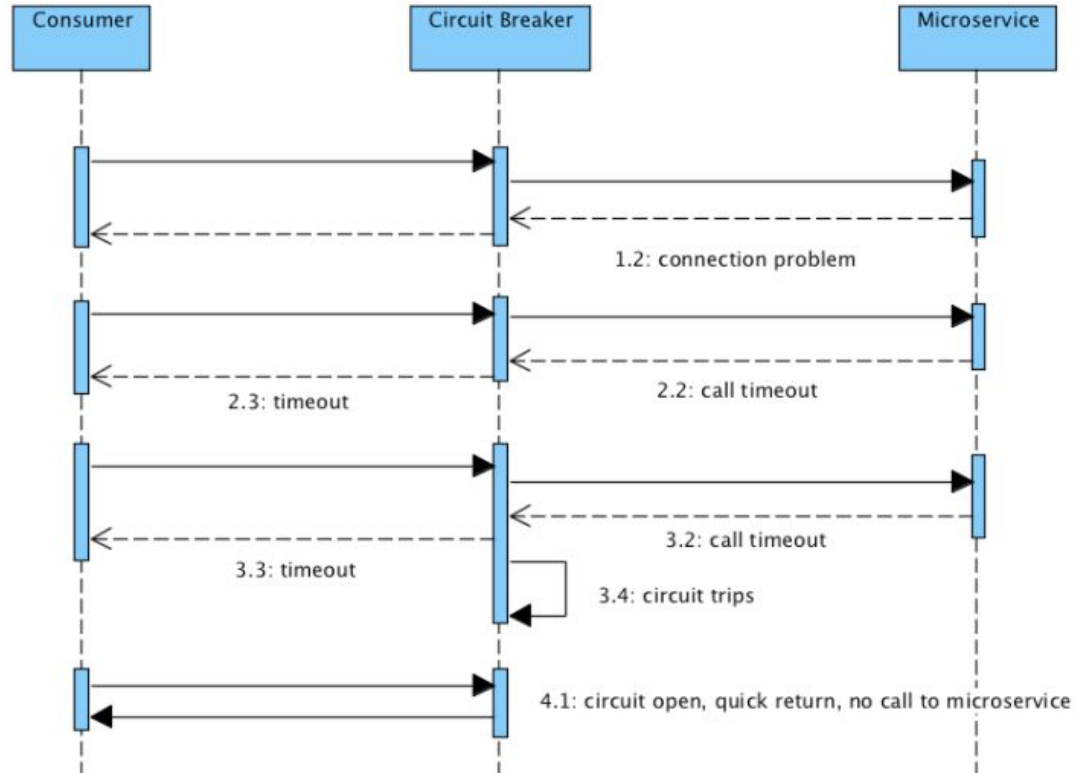
Coding time



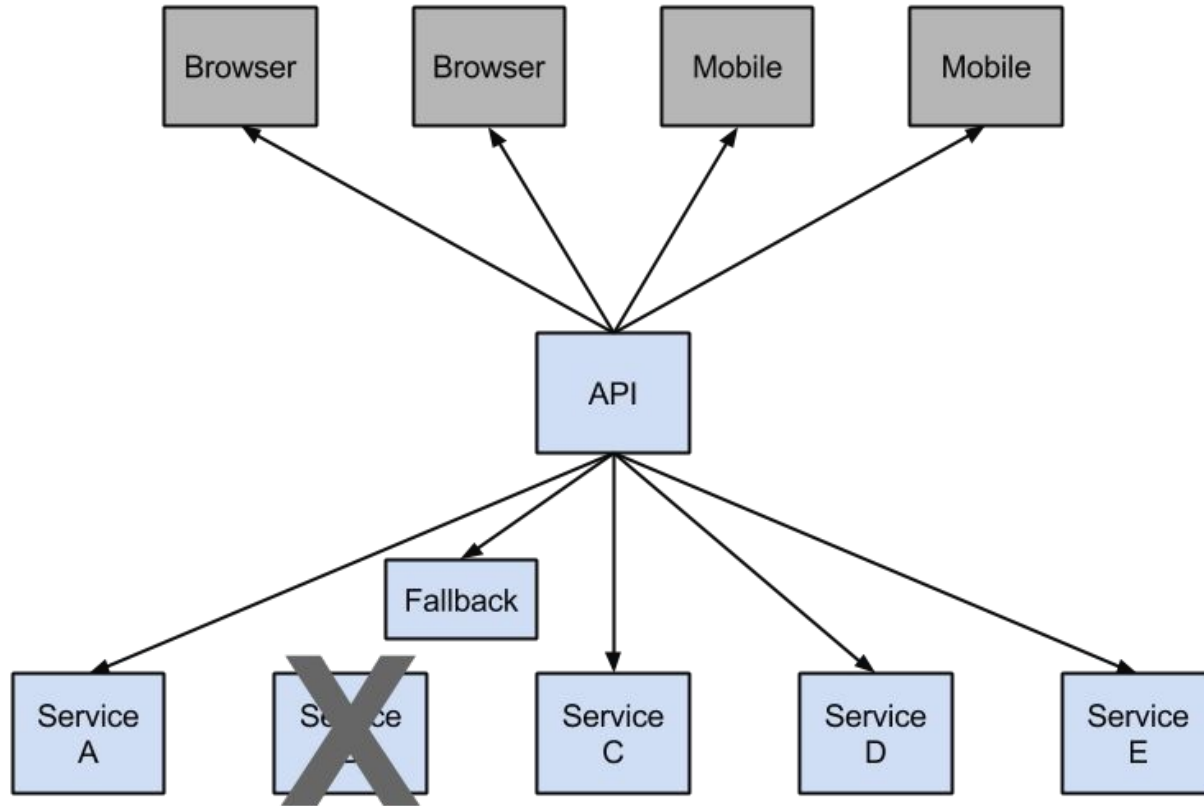
Fault-Tolerance

Distributed Systems

Circuit Breaker



Hystrix Fallback



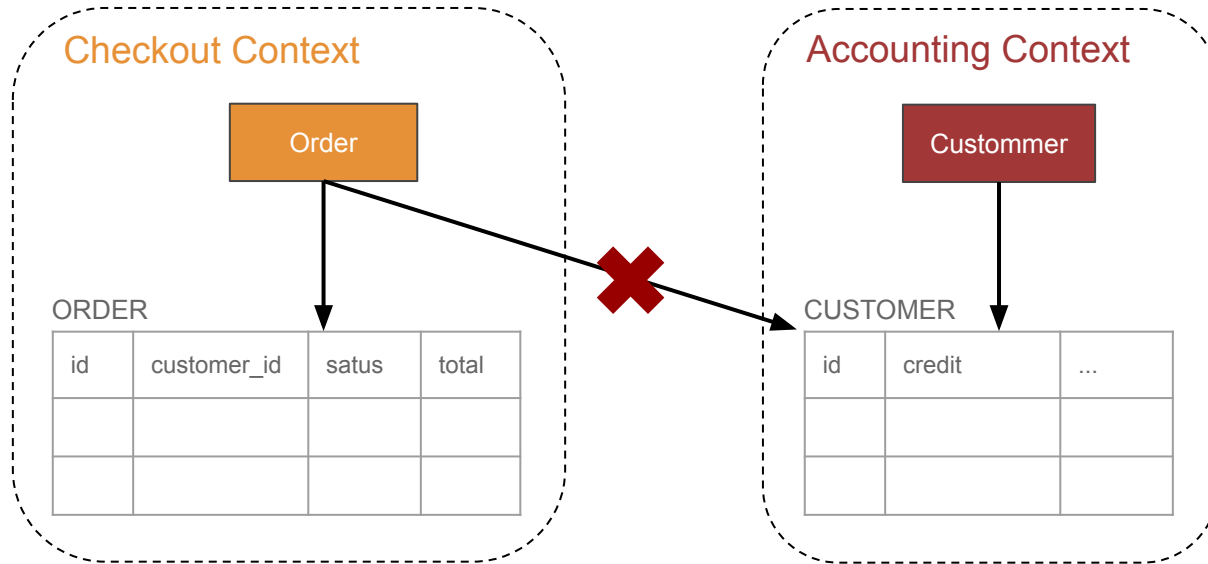
Feign

Declarative REST Client

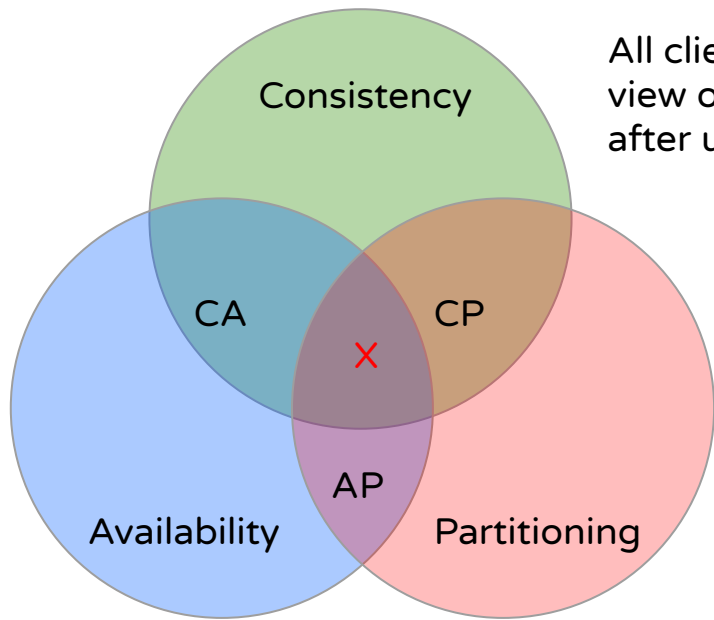
- Discovery Service - Eureka
- Fallbacks - Hystrix
- Request / Response compression

Event-Driven data management

Problem of Distributed Data



CAP theorem

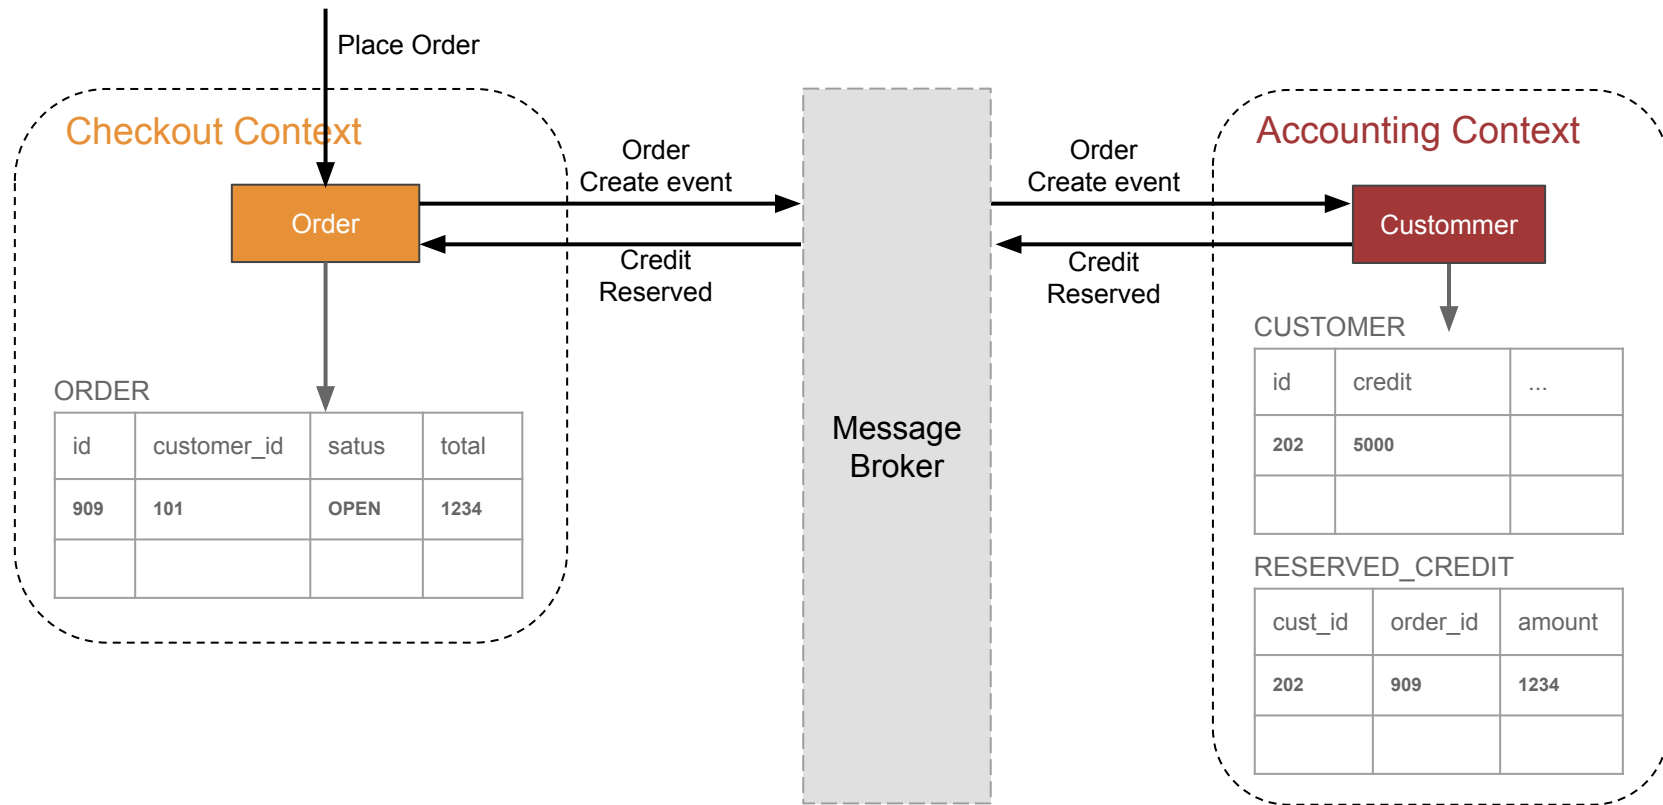


All clients see the same view of data, even right after update or delete

All clients can find a replica of data, even in case of partial node failures

The system continues to work as expected, even in presence of partial network failure

Event-Driven Architecture



Kafka

Message Broker

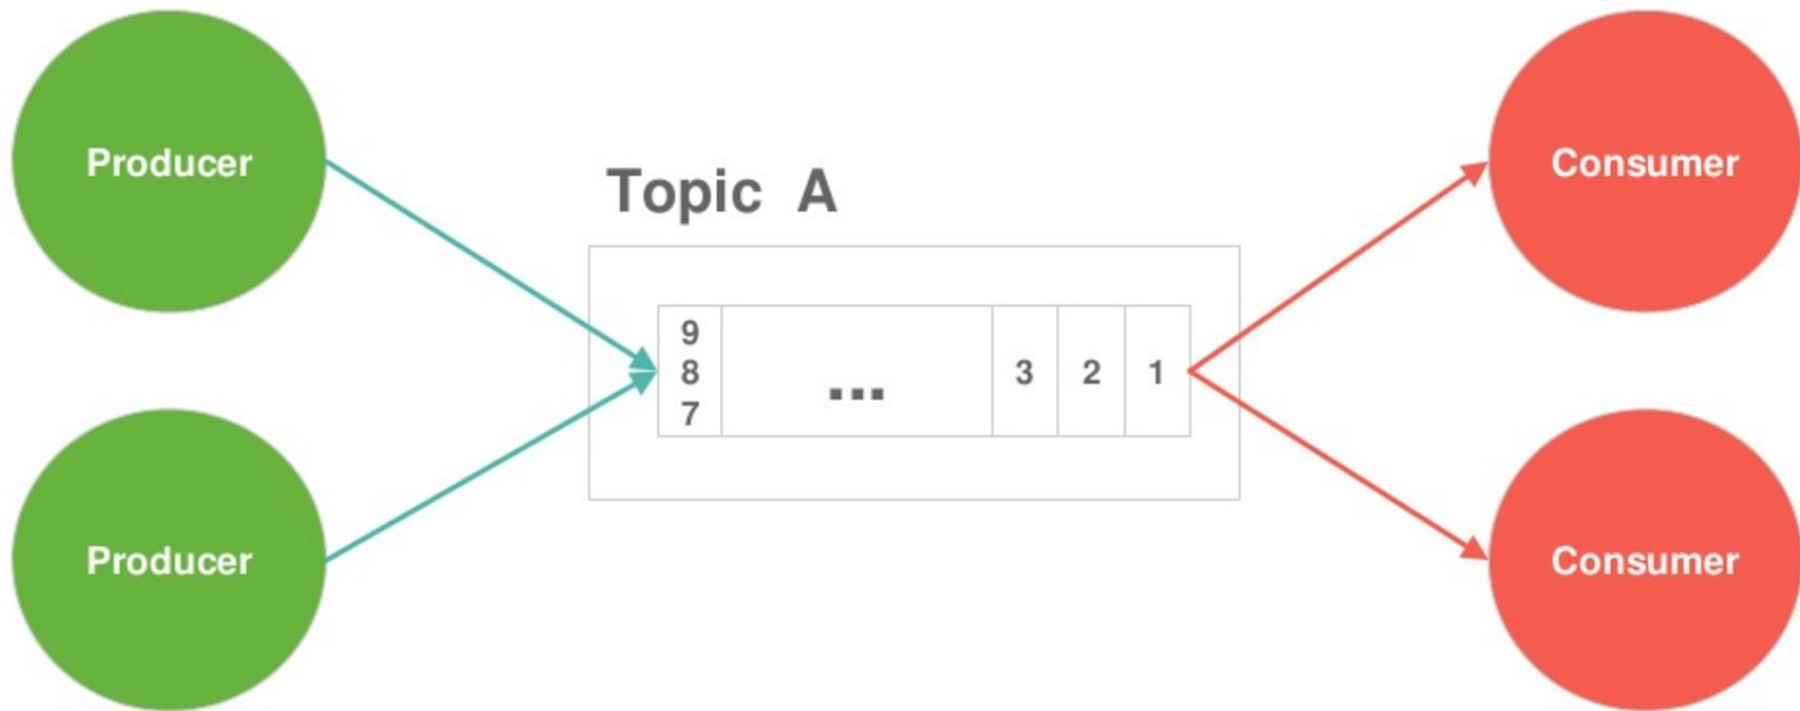
- **Topic** - category or feed name in which messages are published
- **Broker** - server process. The topics live in the broker processes
- **Producer** - publishes data to topics
- **Consumer** - processes subscribed to topic and process the feed of published messages
- **ZooKeeper** - coordinator between the brokers and consumers

Connection the Microservices to Kafka

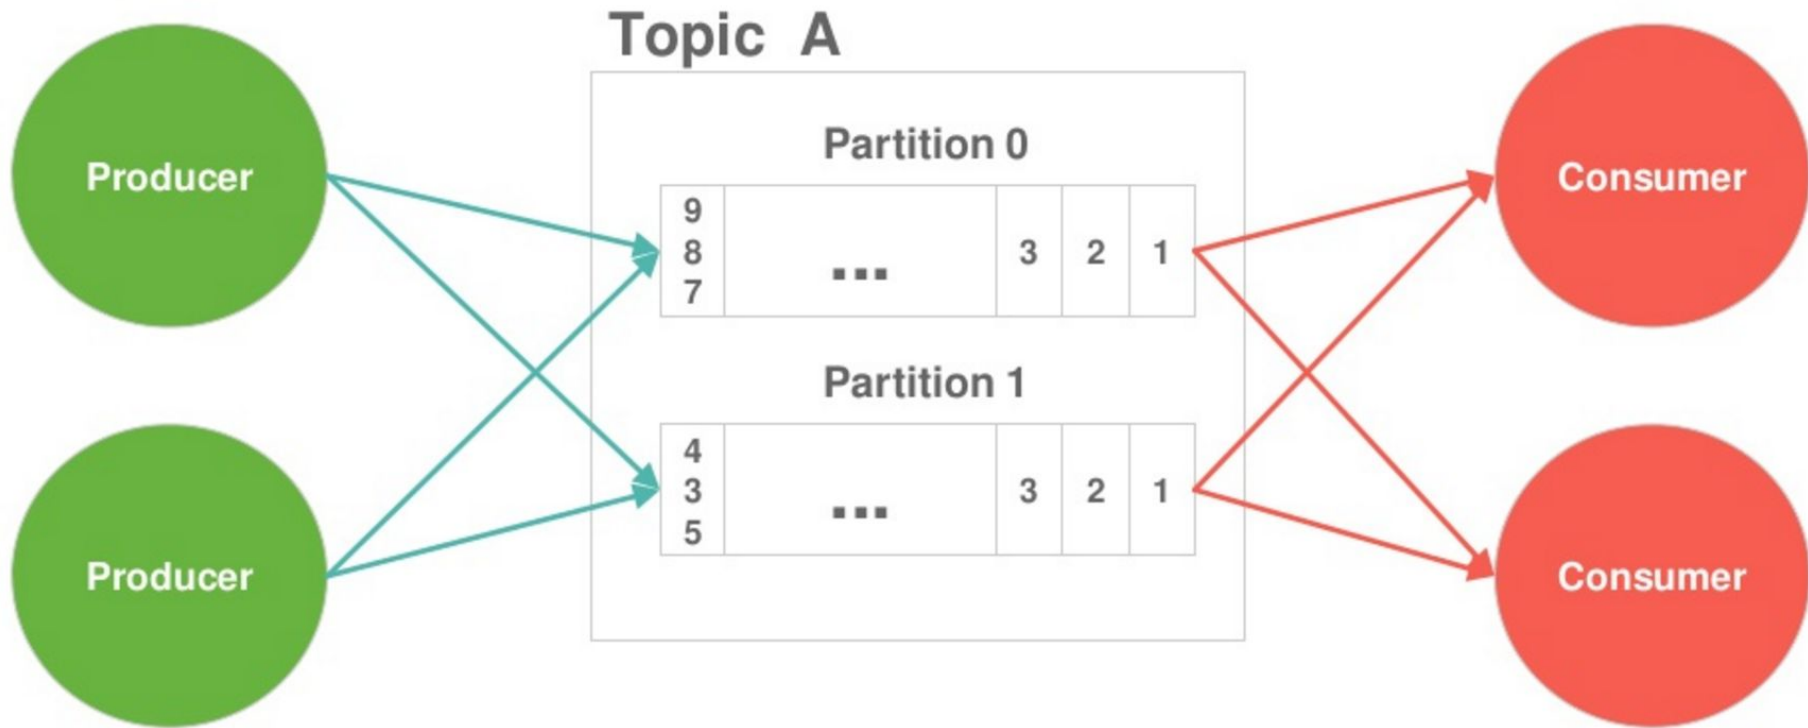
Microservices - Kafka



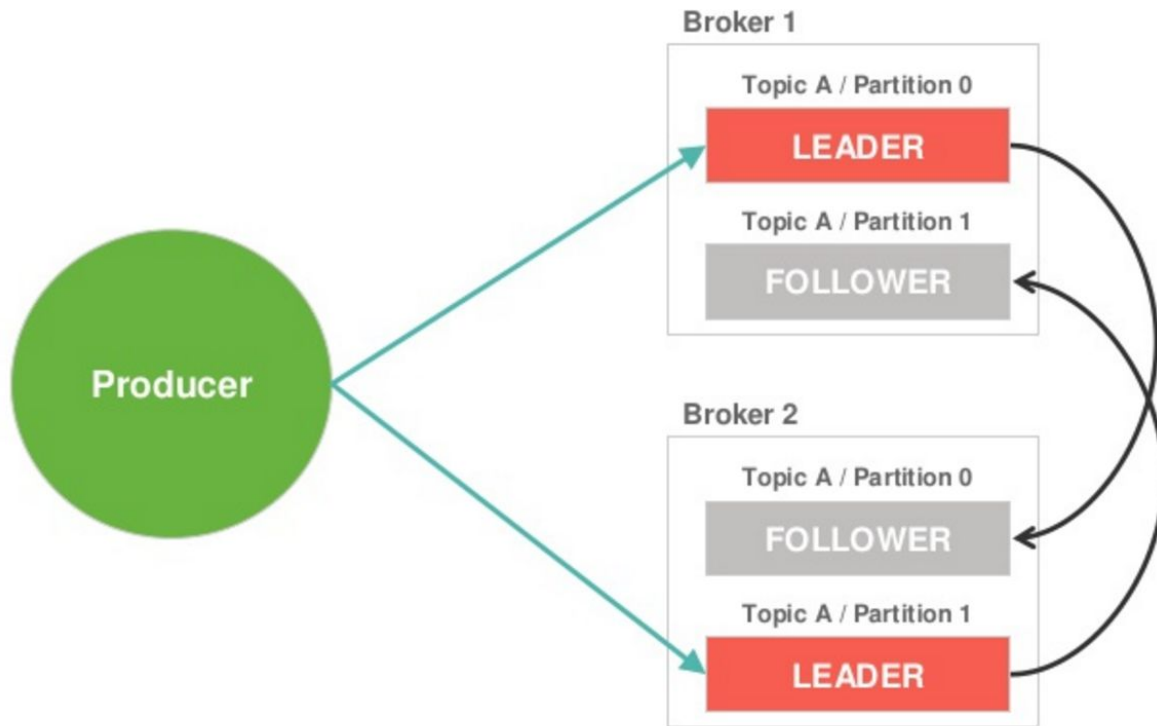
Connection the Microservices to Kafka



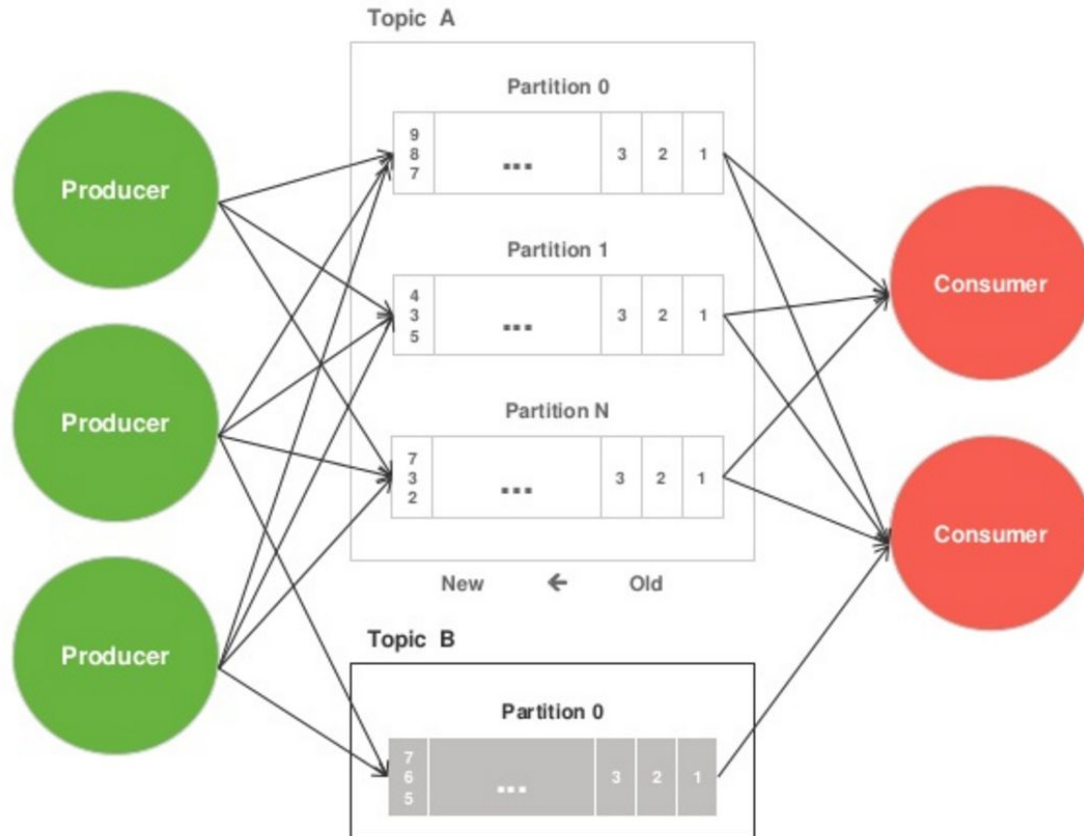
Connection the Microservices to Kafka



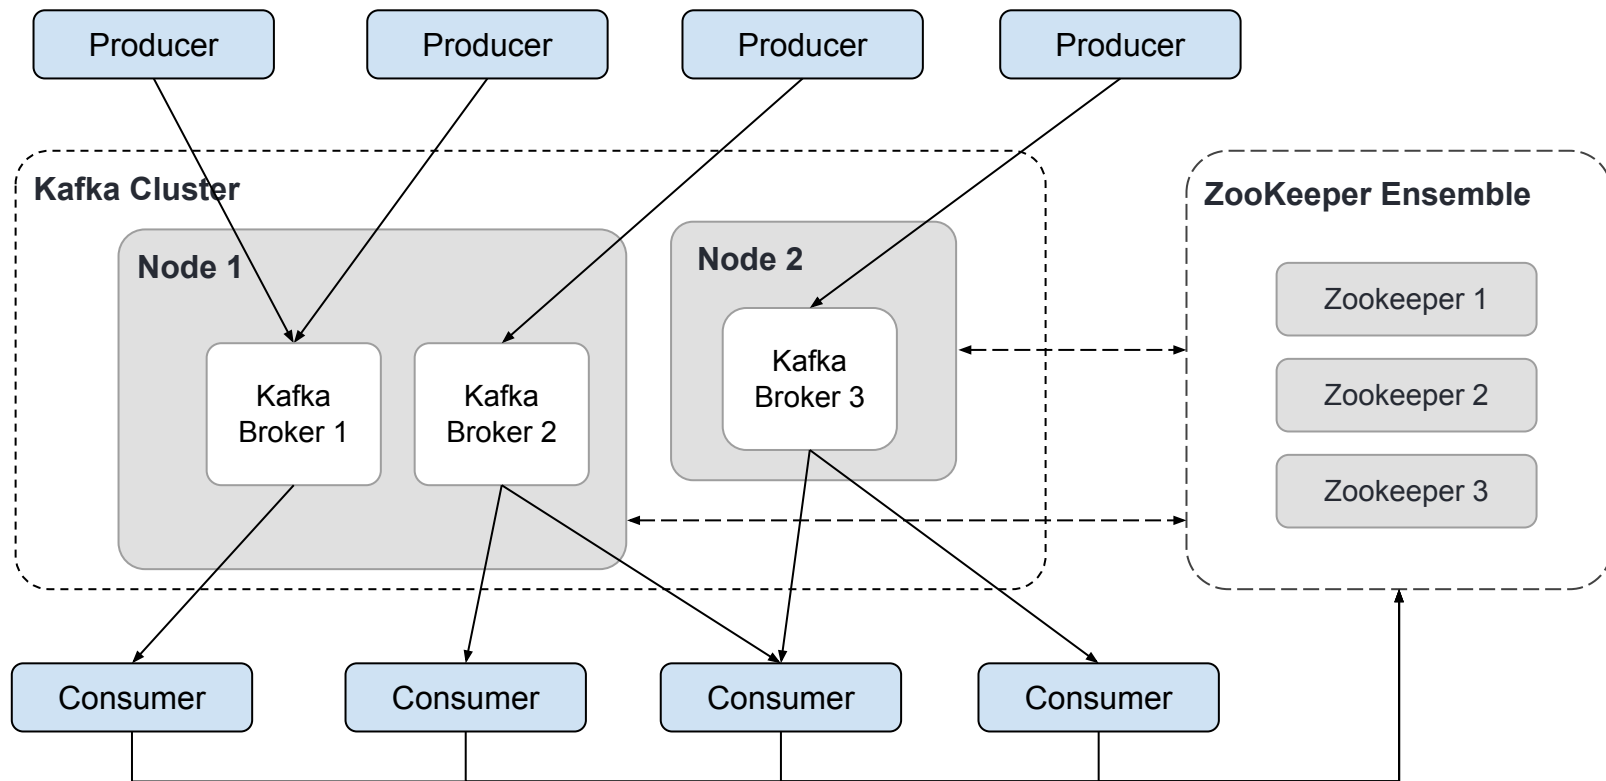
Connection the Microservices to Kafka



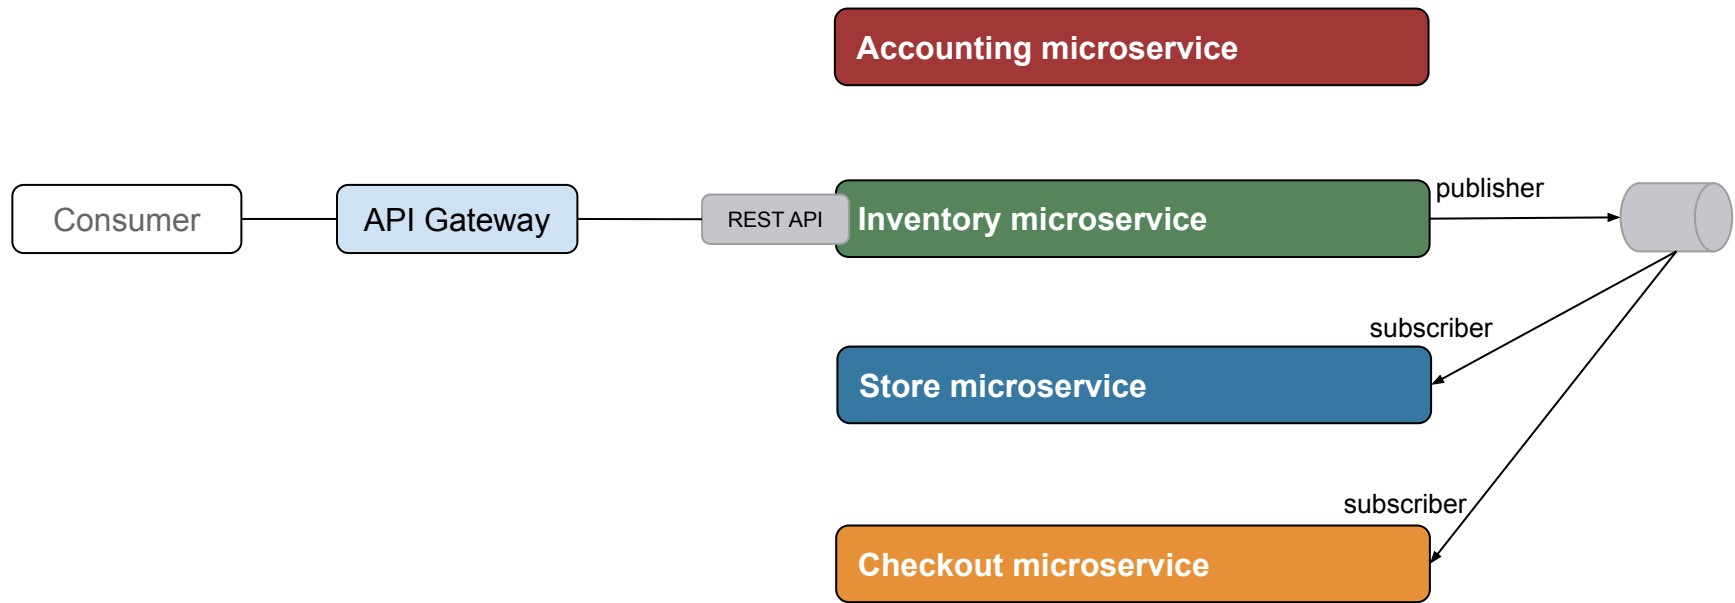
Connection the Microservices to Kafka



Kafka Cluster



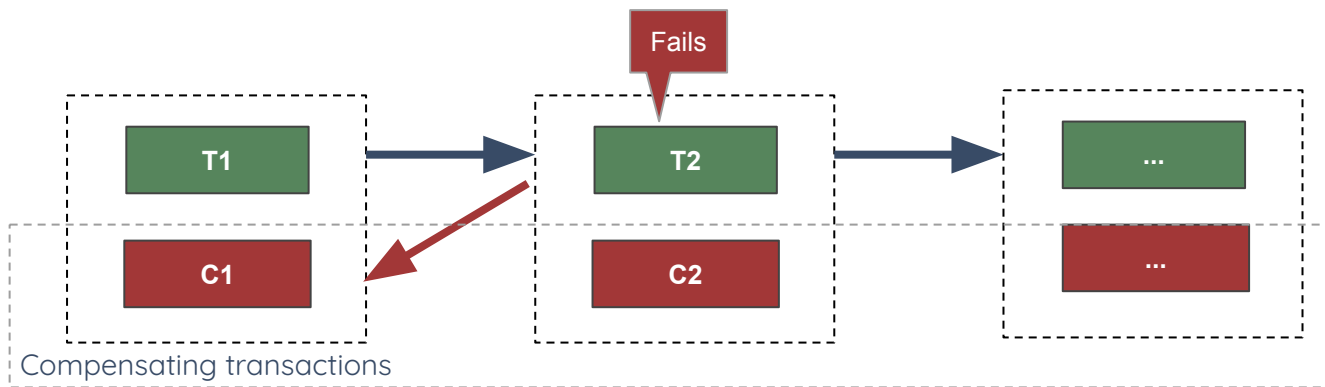
Coding time



Saga

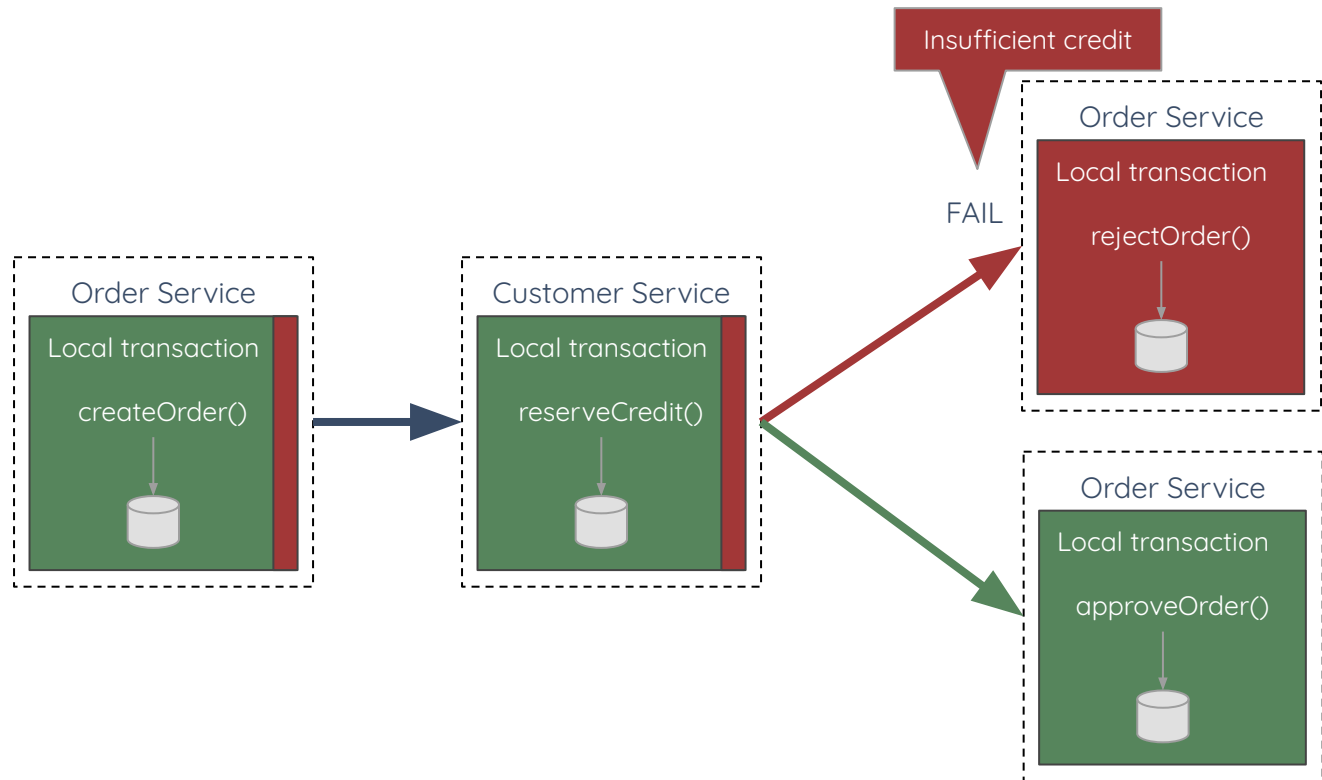
Transaction model

Transaction and Compensation

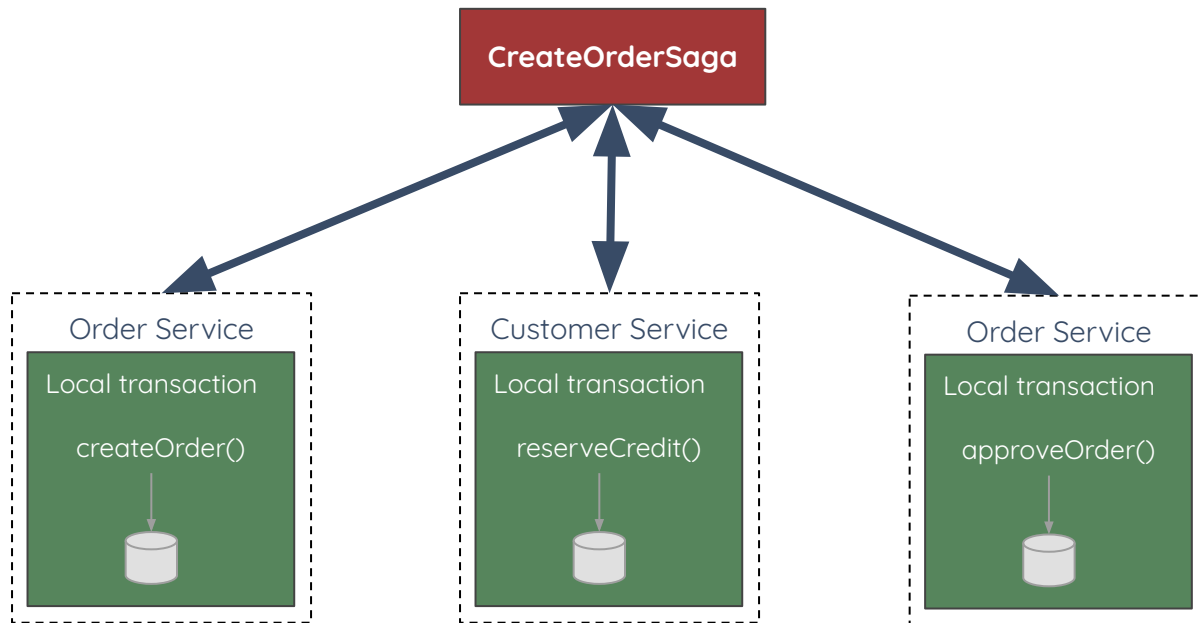


T1 -> T2 -> C1

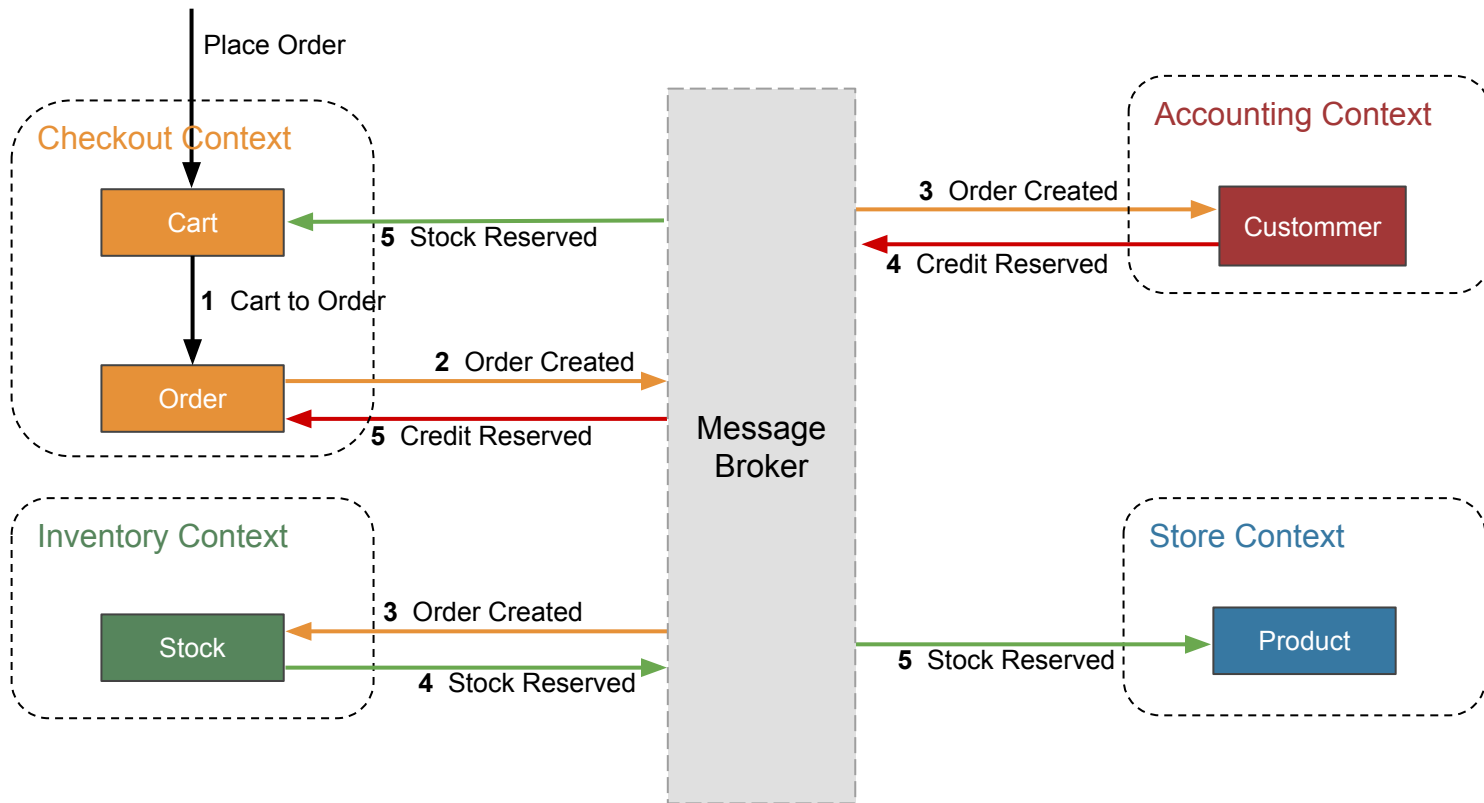
Choreography Coordination



Orchestration Coordination



Coding time



Security

Patterns and Best Practices

Security Model

Availability is guarantee of reliable access to the information by authorized people.

Integrity is the assurance that the information is trustworthy and accurate.

Confidentiality is set of rules that limits access to information.



API Gateway / Perimeter security

- Requests are authenticated and authorized by the gateway
- The public LB cannot send request to apps directly
- Apps trust all traffic they receive by assumption

Pros

- Network setup can virtually guarantee assumptions
- Apps have stateless security

Cons

- Does nothing for internal threats

Basic + Central Auth DB

- All apps get to do authentication and authorization themselves
- Credentials are verified against a central DB
- Basic credentials are passed along in every request

Pros

- Stateless - authenticate every time
- Central user store

Cons

- Auth DB is hit every request
- DB lookup logic needs to be implemented everywhere
- User's credentials can unlock all functionality

Sessions everywhere

- Same as before but each app gets to maintain a session with the client device

Pros

- Auth DB is hit once per session

Cons

- Hard to manage all the sessions
- No single sign on
- DB lookup logic needs to be implemented everywhere
- User's credentials can unlock all functionality

API Tokens

- Username and password is exchanged for a token at a centralized auth server
- Apps validate the token for each request by hitting the auth server

Pros

- App don't see user credentials

Cons

- Auth server bottleneck
- Token provides all or nothing access

Microservices

Security Concerns

- Central user store bottleneck
- Single Sign On
- Statelessness
- User credentials
- Fine grained authorization
- Interoperability with non browser

JWT Token

eyJhbGciOiJSUzI1NiIsInR5cCI6IkpXVCJ9.eyJ1c2VyX25hbWUiOiJyZWFKZXkiLCJzY29wZSI6WyJyZWFKiiwid3JpdGUxSwicm9sZXMl0lt7ImF1dGhvcm10eSI6IlJPTEVfVVNFUiJ9XSwiZnVsbg5hbWUiOiJkdyxpbiB0aWNvYXJhIiwiaXhwIjoxNTg5OTI1OTA3LCJhdXRob3JpdGlscyI6WyJST0xFX1VTRVIiXSwianRpIjoieYzA1YzI4YjctODBmOC00ZTY2LTlkMzktNzFjYmIyODY1ZWE0IiwiaY2xpZW50X2lkIjoieY2xpZW50SWQifQ.pahqf0yYUD8BFBumvULGGEgu99wMWG4UWHlkgrrT0zCEPvnjTKHKcE_2oGtKjrtxKPgFL8ESM5ur8JB_kCL0pWB1BLQ3ZE2t8-rJOGNvpbxflbCh8CxrYbdTe75hmDJb3g-BUxdxxg64lfqD0bGC0eILfJf1Djoe1BY3uCg1gs7dt6lthT4GhTb4pmq6j70z70ztIHqn6bwnHFjJwwnxyZ43uwnc5Er1snKkUmgg3xWq_9R82iHQ06eSOIWbdRtPWbhRSKi-70aG8wSoun-0bVm9aS0ayq1CZmnHDNrIhRdgXHrnavPBHBX9Ye1Tv2_rILcyHHFYnDilG1PISlenjq

HEADER: ALGORITHM & TOKEN TYPE

```
{
  "alg": "RS256",
  "typ": "JWT"
}
```

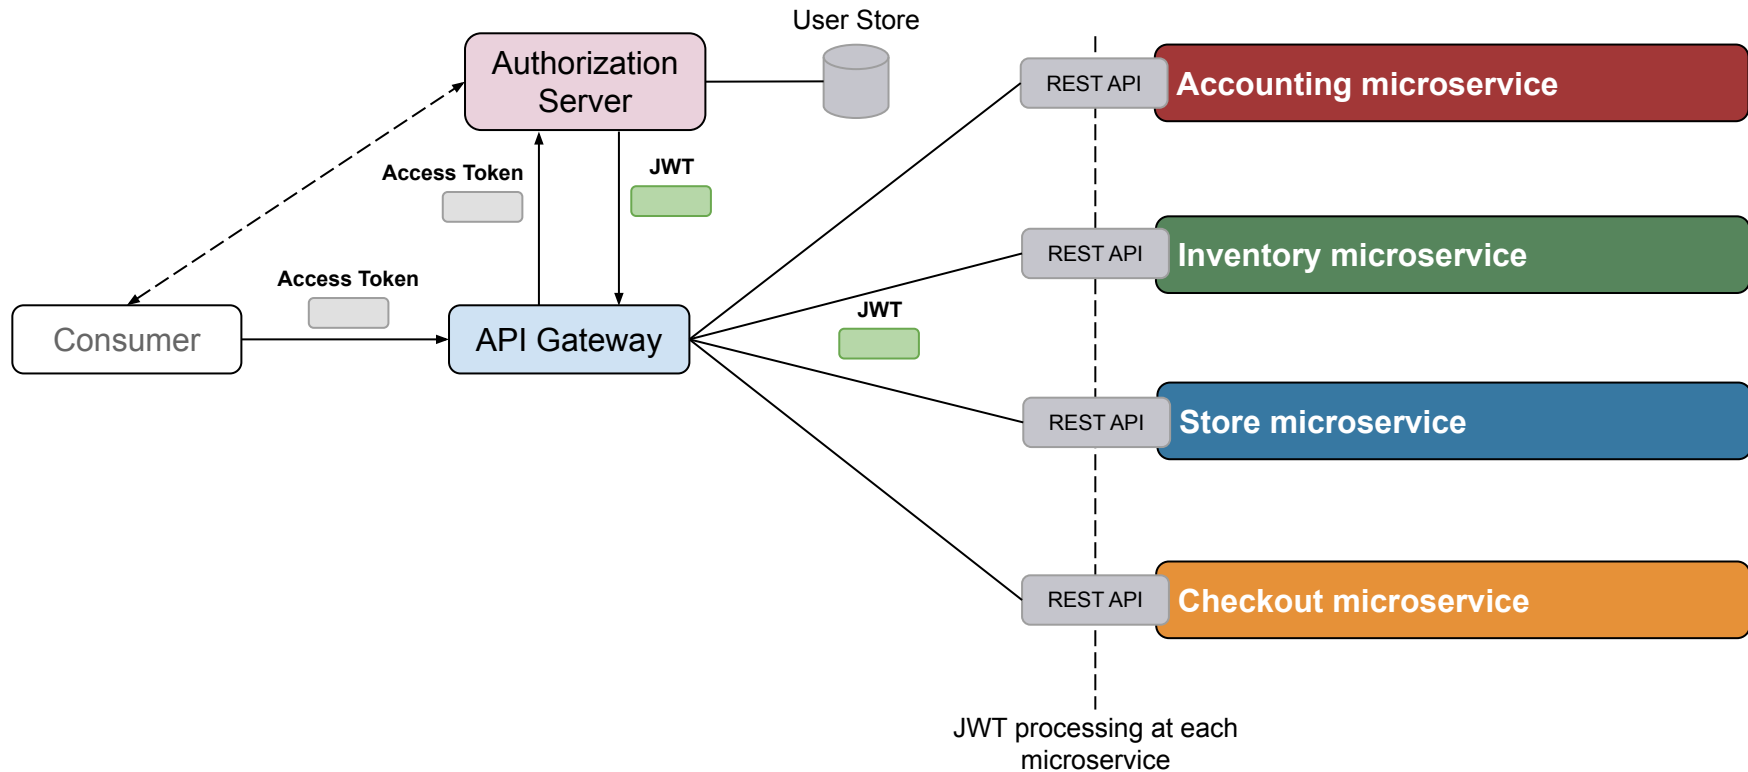
PAYLOAD: DATA

```
{
  "user_name": "reader",
  "scope": [
    "read",
    "write"
  ],
  "roles": [
    {
      "authority": "ROLE_USER"
    }
  ],
  "fullname": "Calin Nicoara",
  "exp": 1589925907,
  "authorities": [
    "ROLE_USER"
  ],
  "jti": "c05c28b7-80f8-4e66-8d39-71cbb2865ea4",
  "client_id": "clientId"
}
```

Oauth2 Concepts

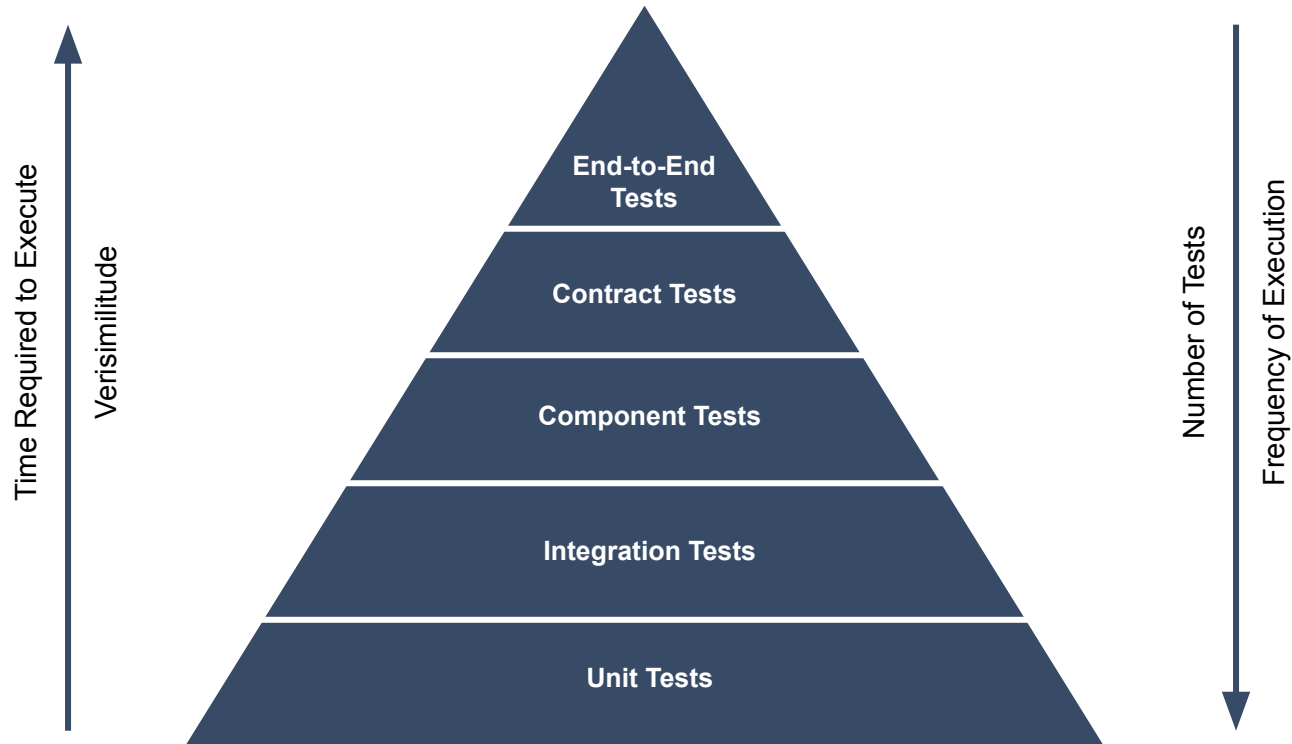
- **Resource Owner** - The user who authorizes an application to access his account. The access is limited to the scope.
- **Resource Server** - A server that handles authenticated requests after the **client** has obtained an **access token**
- **Client** - An application that access protected resources on behalf of the resource owner.
- **Authorization Server** - A server which issues access tokens after successfully authenticating a **client** and **resource owner** and authorizing the request.
- **Access Token** - A unique token used to access protected resources
- **Scope** - A Permission (Example: READ, WRITE)
- **JWT** - JSON Web Token is a method for representing claims securely between two parties
- **Grant type** - a method of acquiring an access token.

OAuth2 and OpenID Connect



Testing Strategies

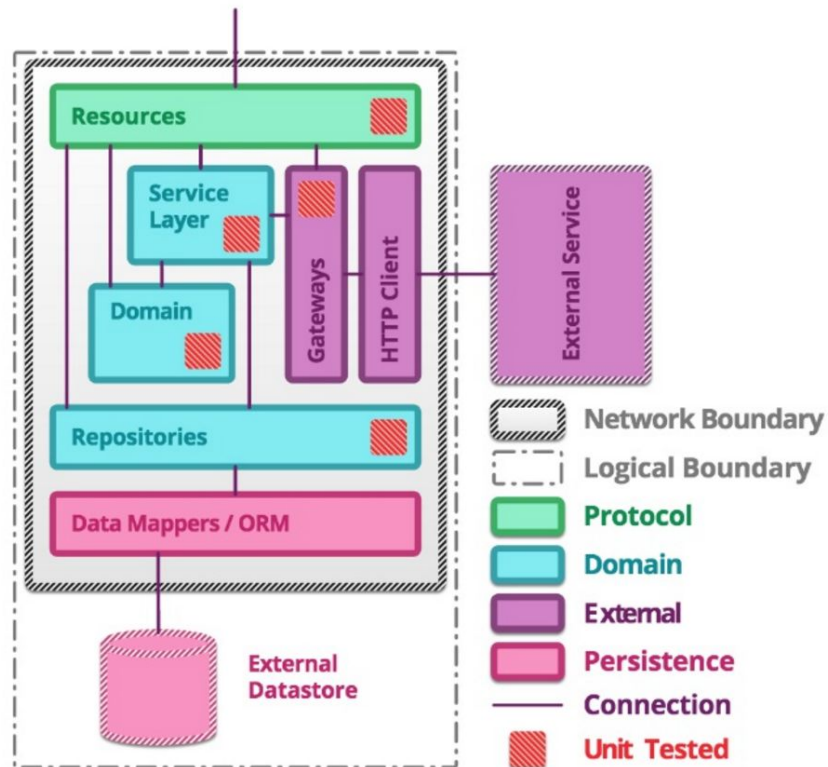
Testing Pyramid



Unit Testing

Make sure that the each component behave as expect.

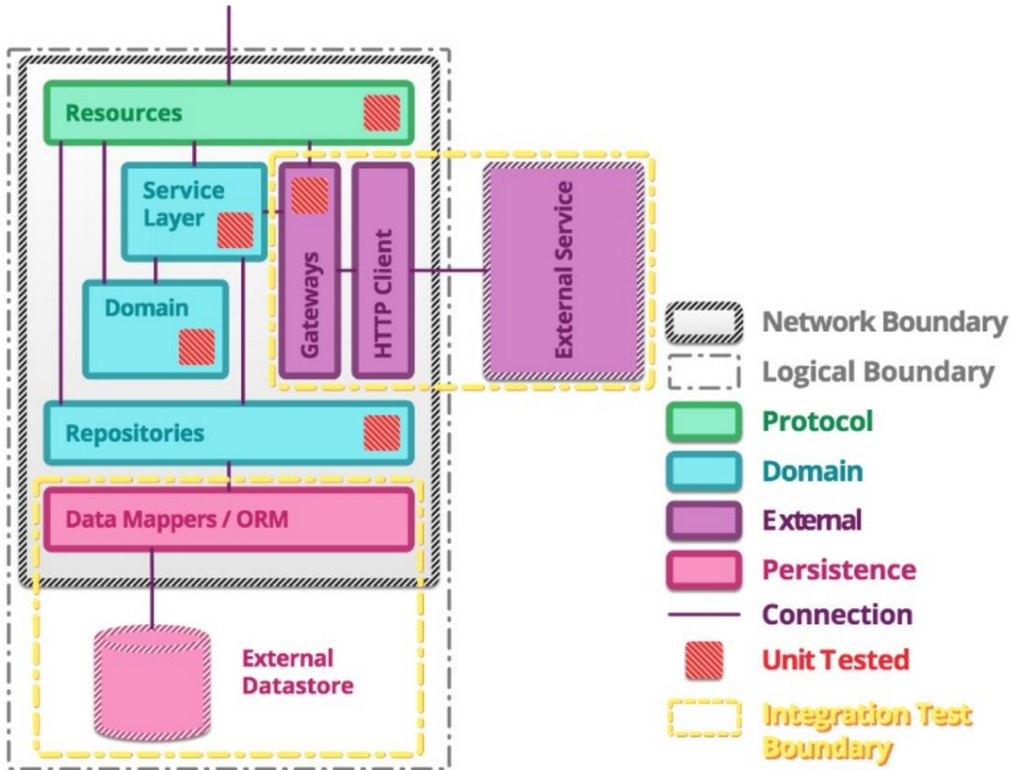
- Domain
- Gateways
- Resources
- Persistence



Integration Testing

Verifying the communications and interactions between components

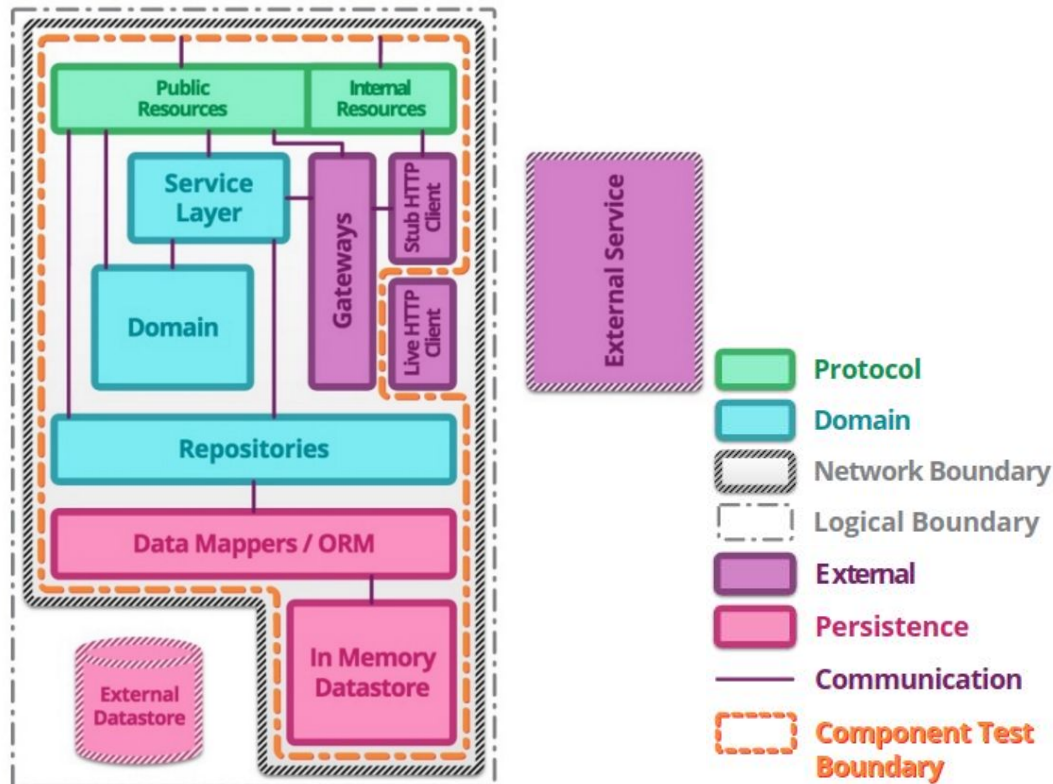
- Persistence
- Gateways



Component Testing

Testing the service in isolation

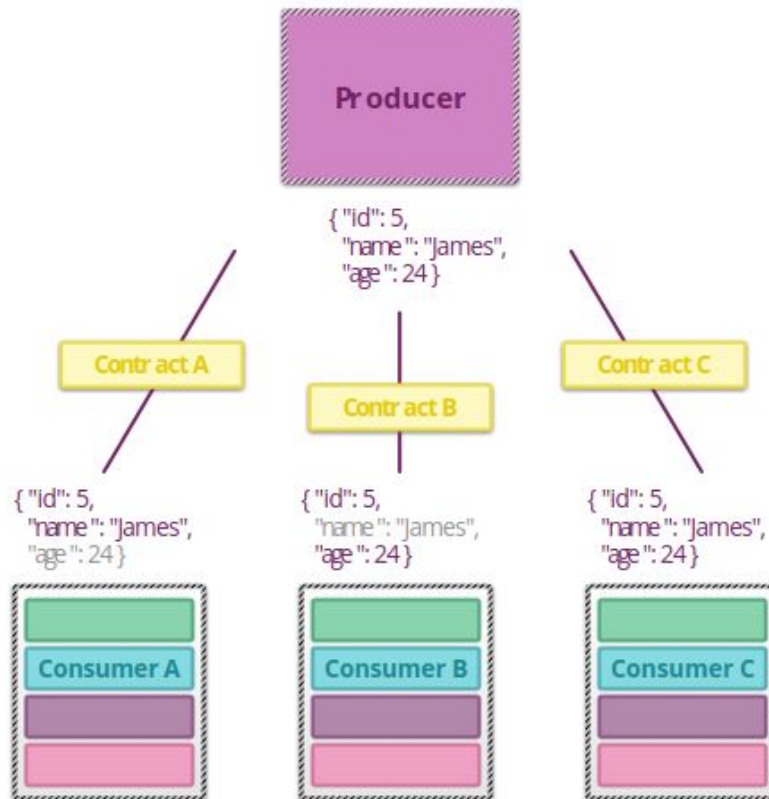
- In Memory DB
- Stub HTTP Client
- Simulate external service



Contract Testing

Verifying the agreements between producer and consumers services

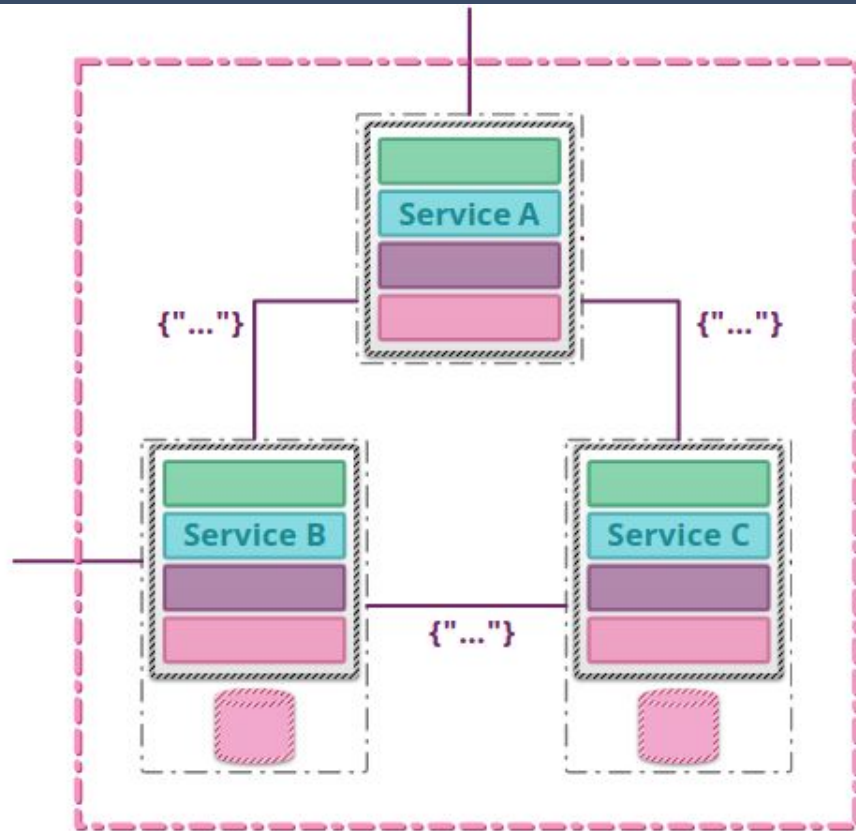
- Input & Output data structure
- Calls attributes



End-to-End Testing

Testing the behaviour of the entire system

- Business flow
- Proxies
- Load Balancers



Deployment Strategy

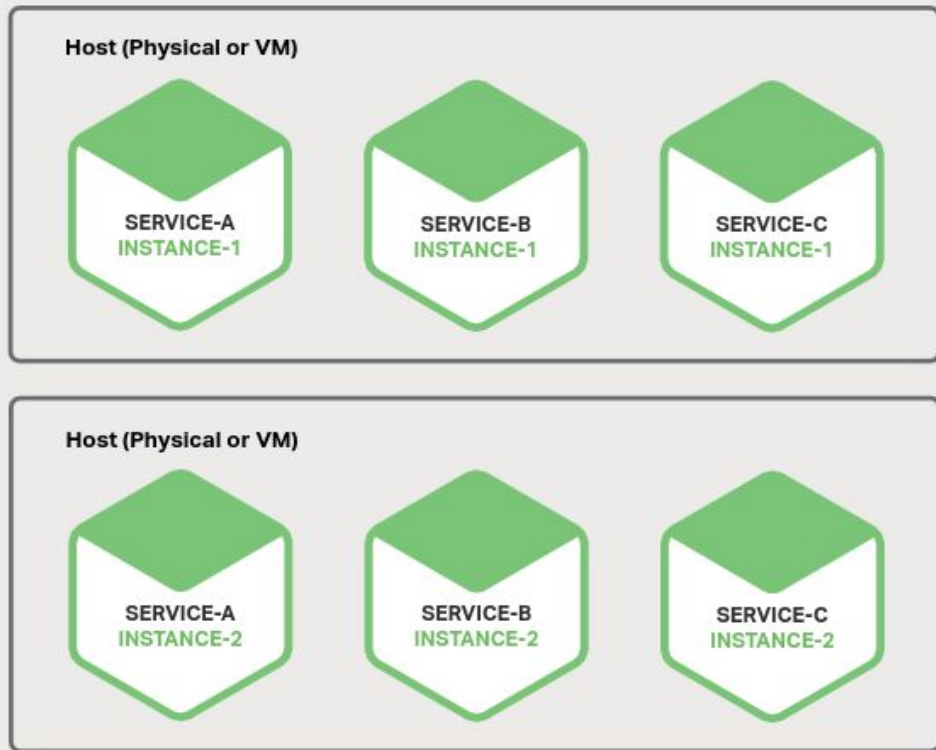
Multiple Service Instances per Host

Pros

- Efficient resource utilization
- Fast deployment

Cons

- Poor isolation
- Difficult to limit resource utilization
- Risk to dependency version conflicts
- Poor encapsulation of implementation technology



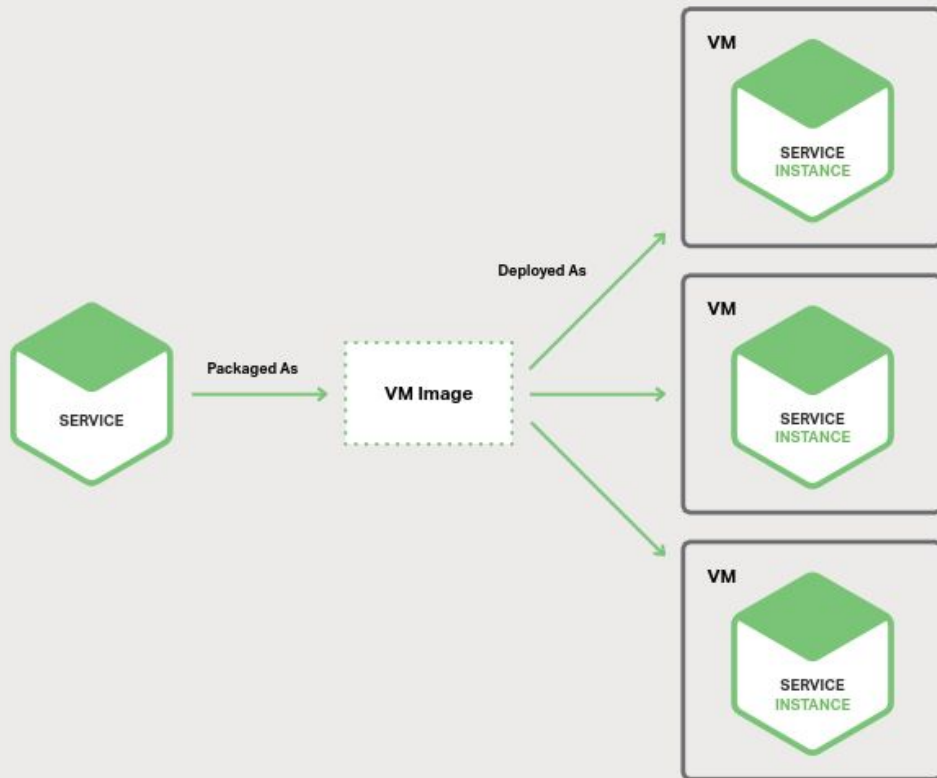
Service Instance per Virtual Machine

Pros

- Great isolation
- Great manageability
- VM encapsulates implementation technology

Cons

- Less efficient resource utilization
- Slow deployment



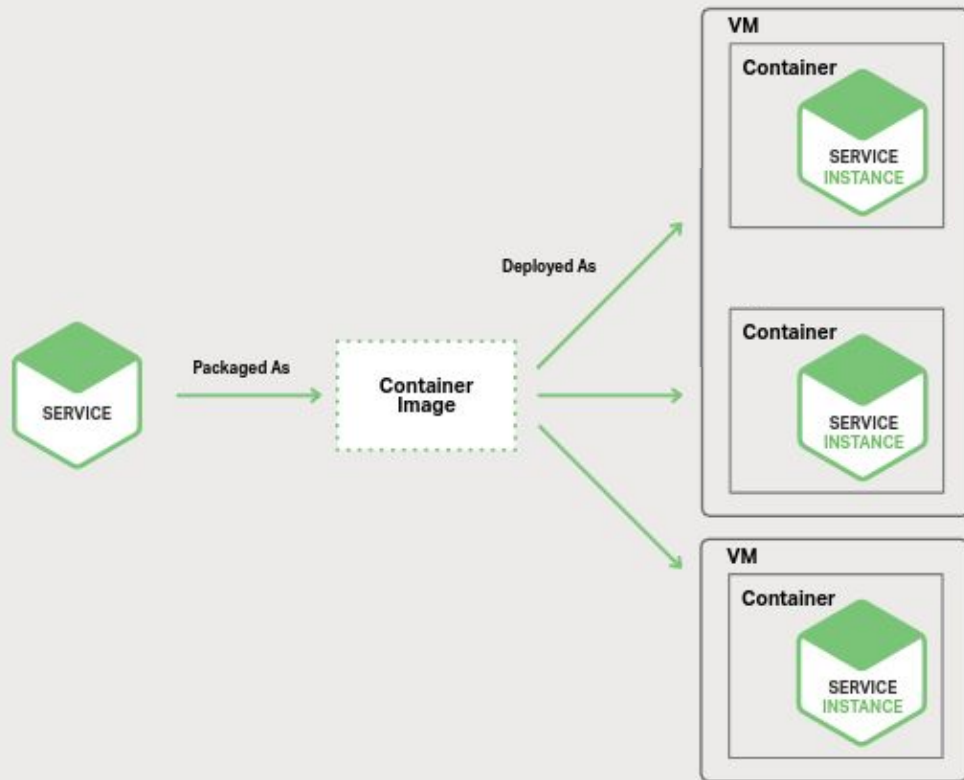
Service Instance per Container

Pros

- Great isolation
- Great manageability
- Container encapsulates implementation technology
- Efficient resource utilization
- Fast deployment

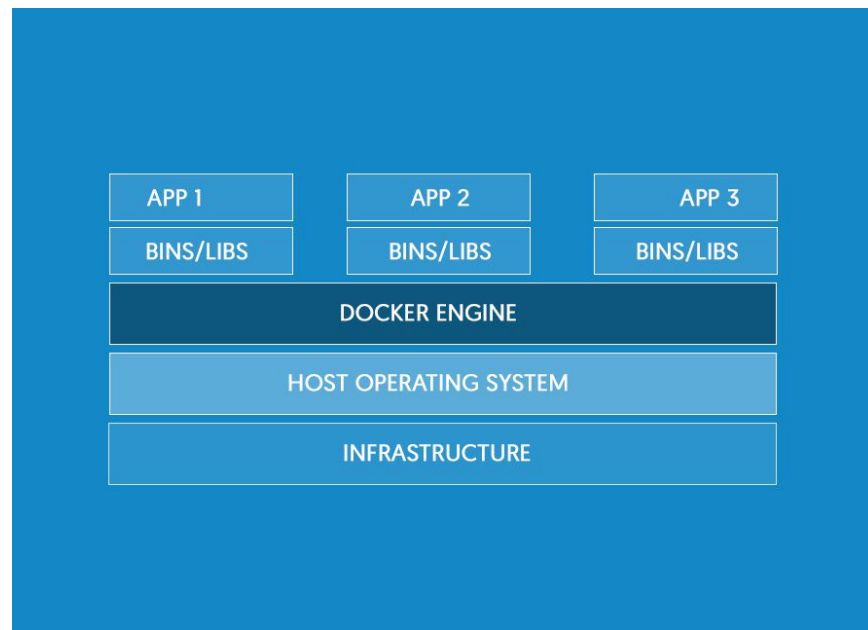
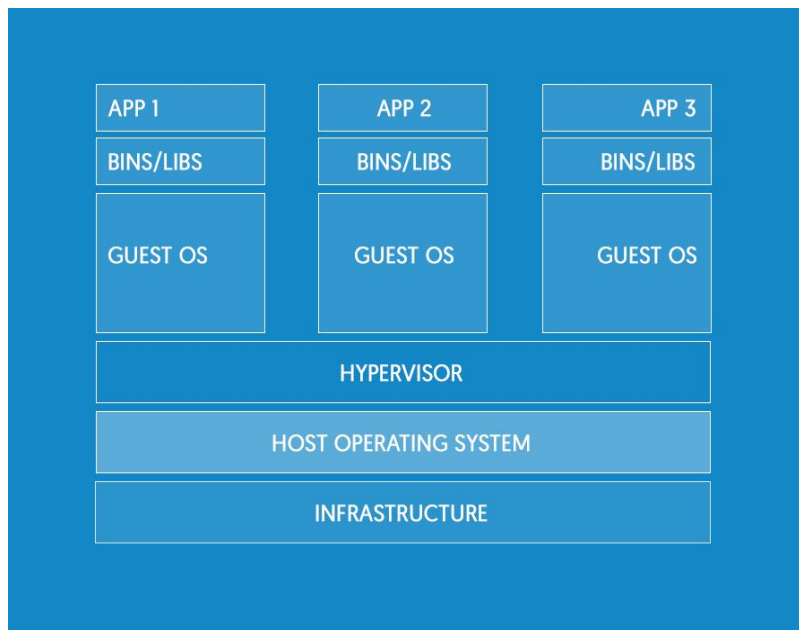
Cons

- Immature infrastructure for deploying containers

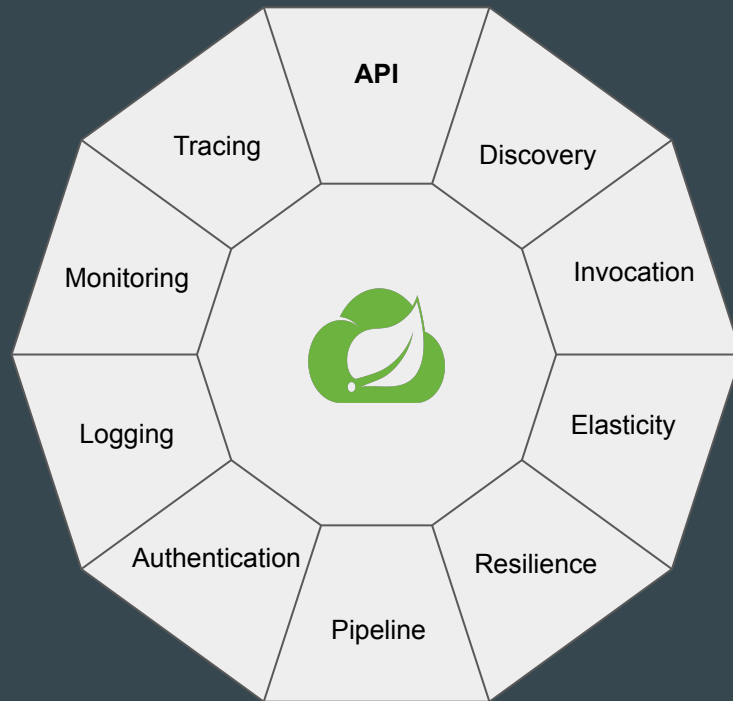


Deployment

With containers



Microservices requirements

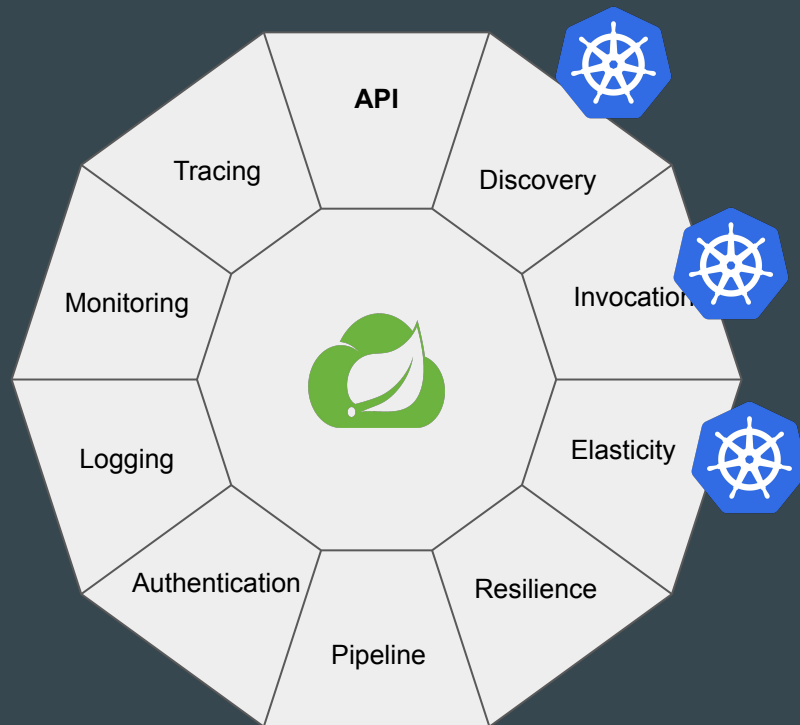


Kubernetes

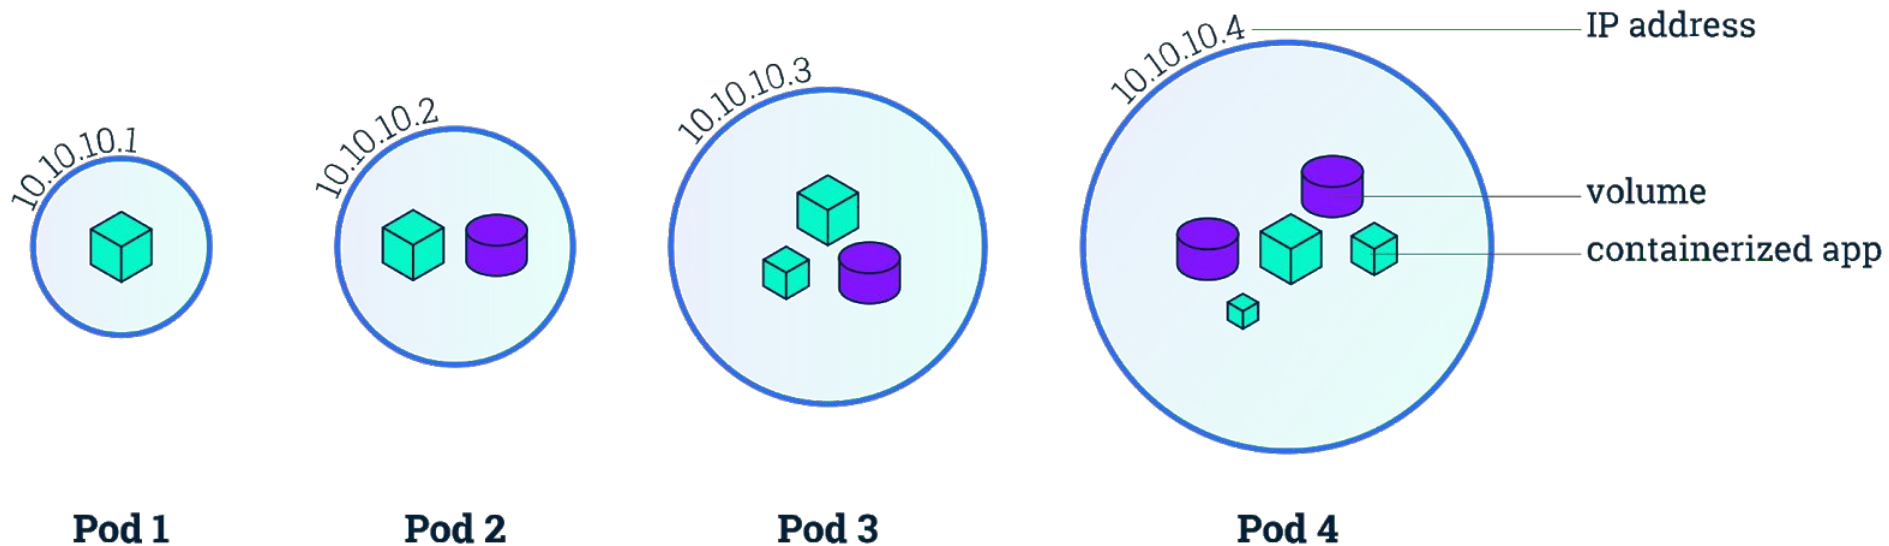


Kubernetes coordinates a highly available cluster of computers that are connected to work as a single unit.

What Kubernetes covers

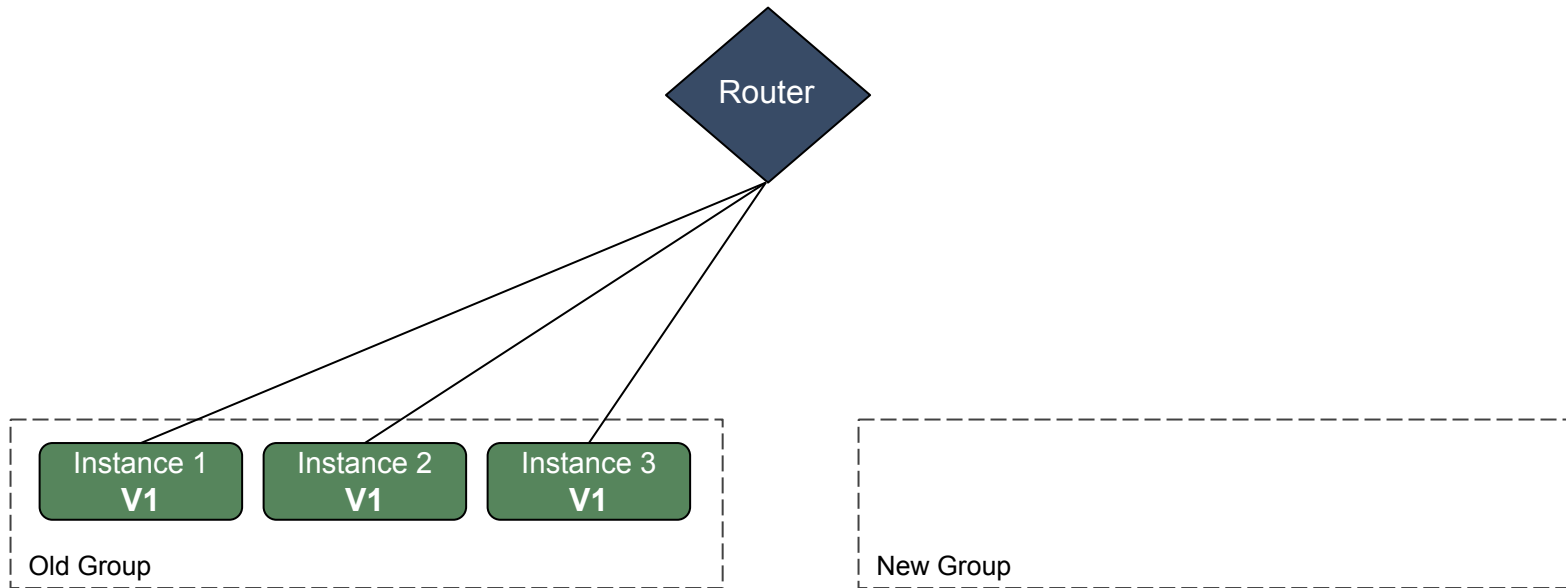


Kubernetes Pods

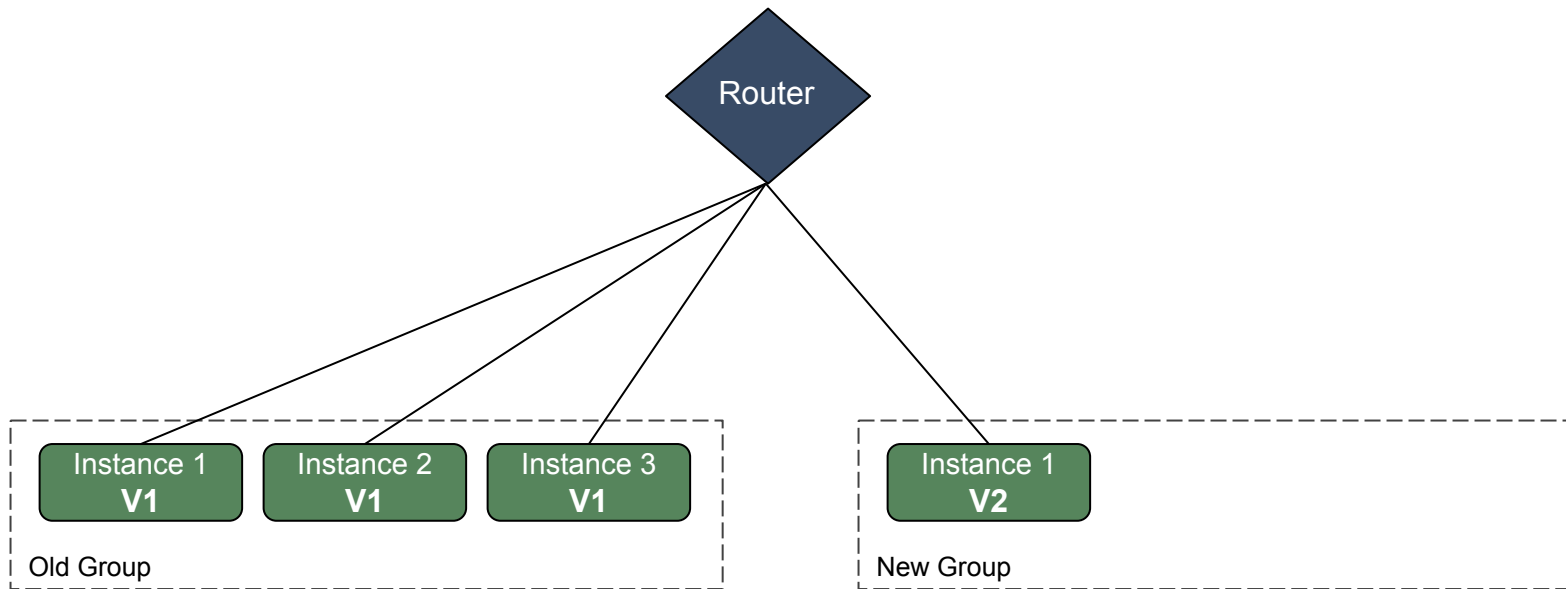


Deployment Patterns

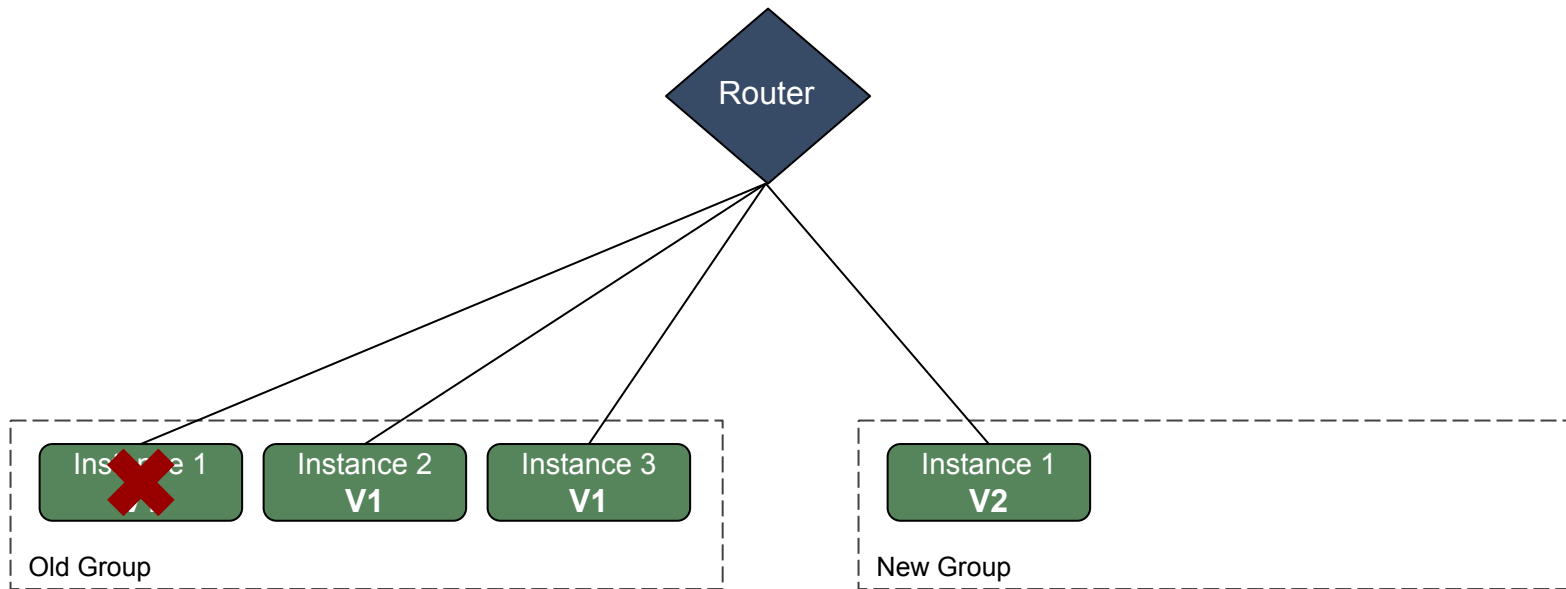
Rolling Update



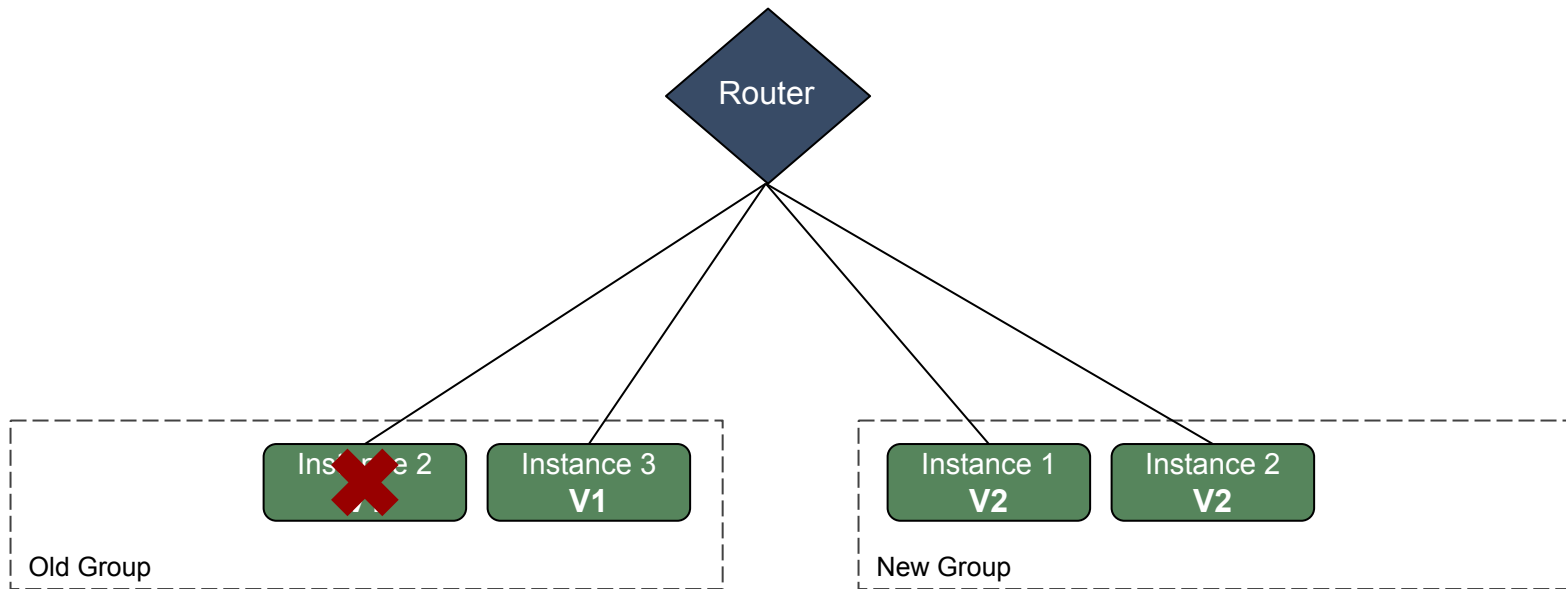
Rolling Update



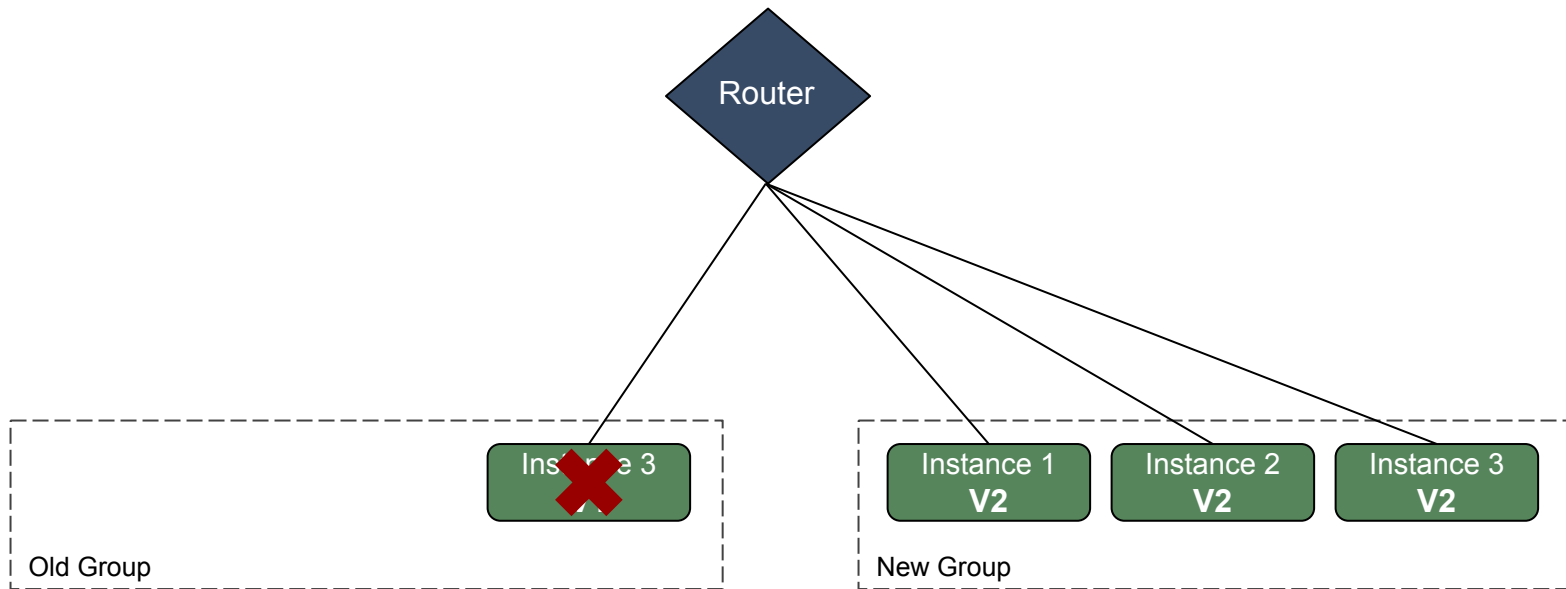
Rolling Update



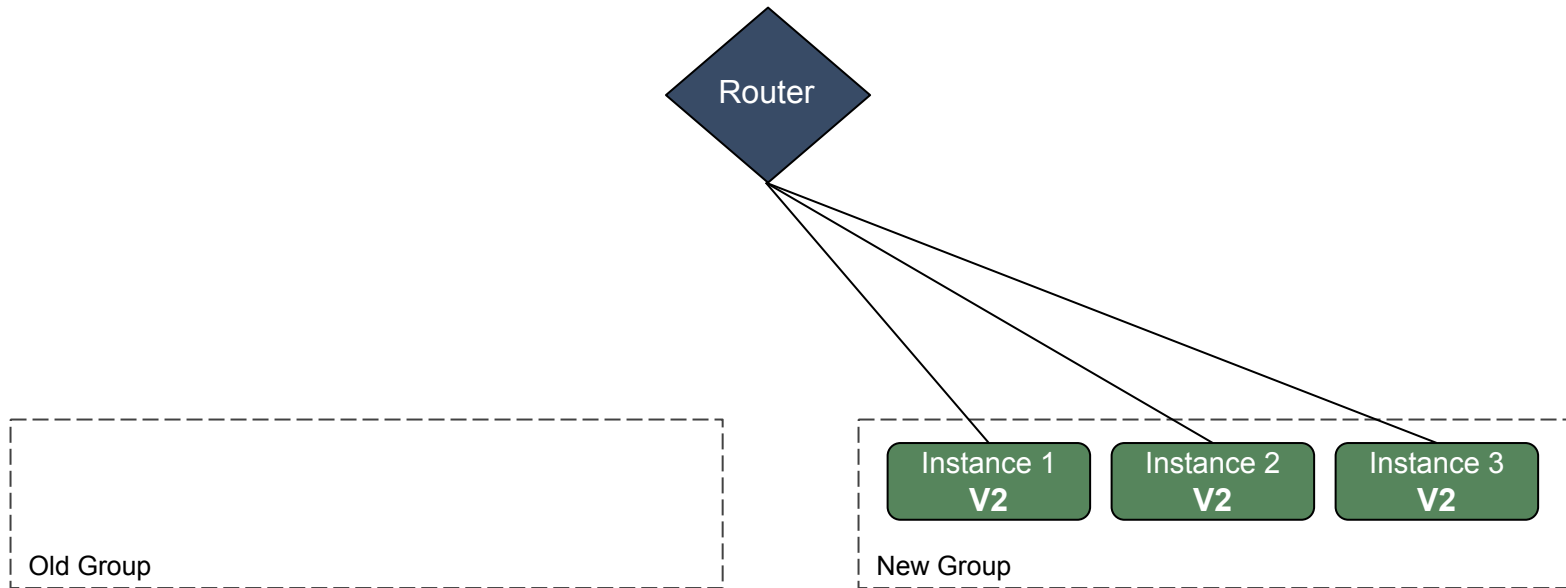
Rolling Update



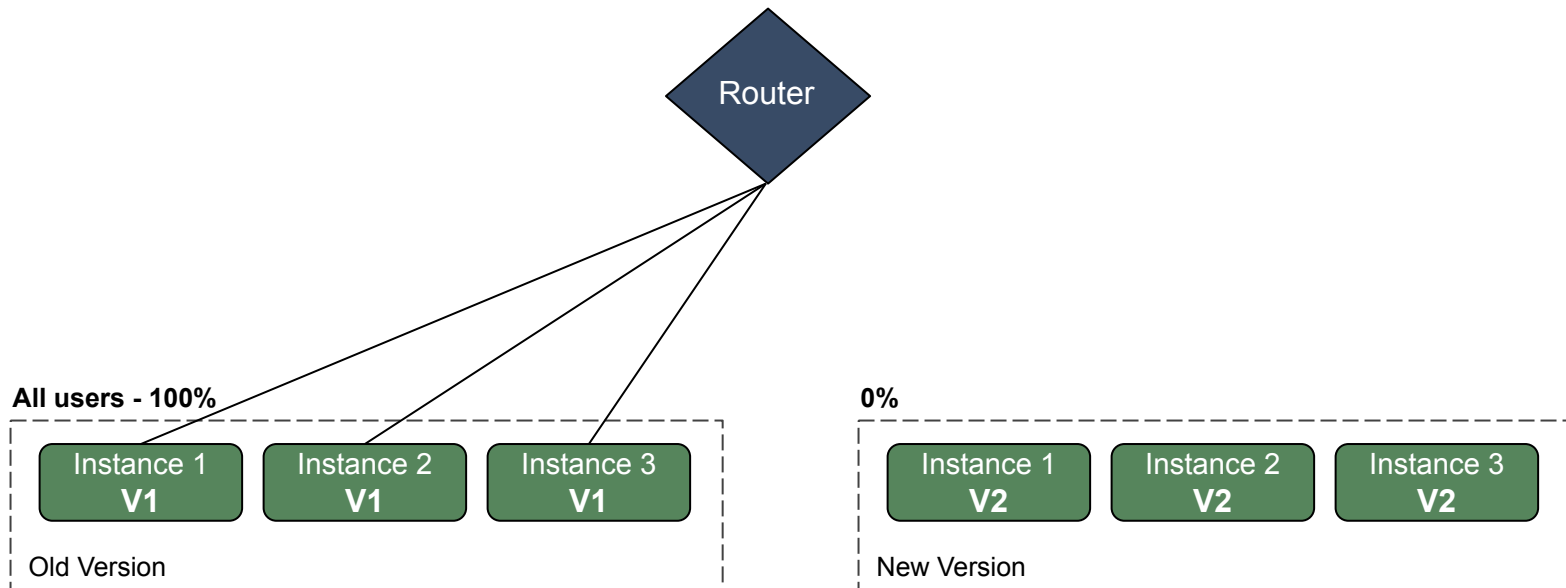
Rolling Update



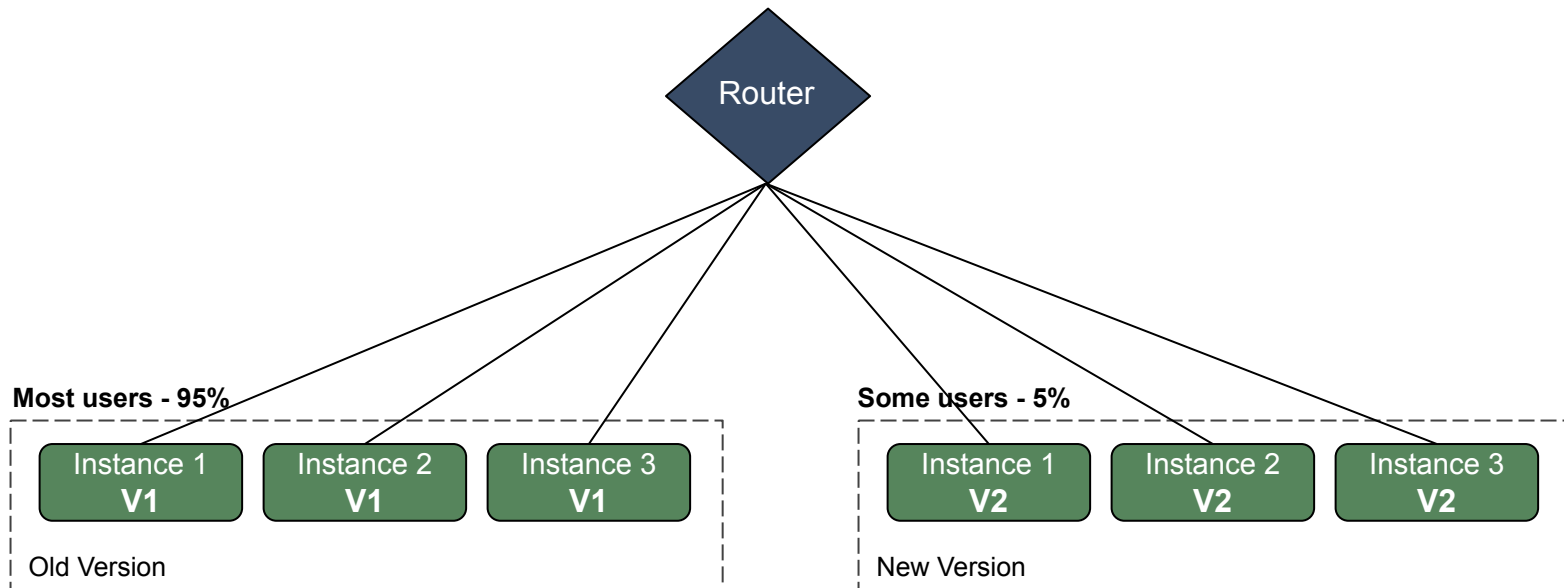
Rolling Update



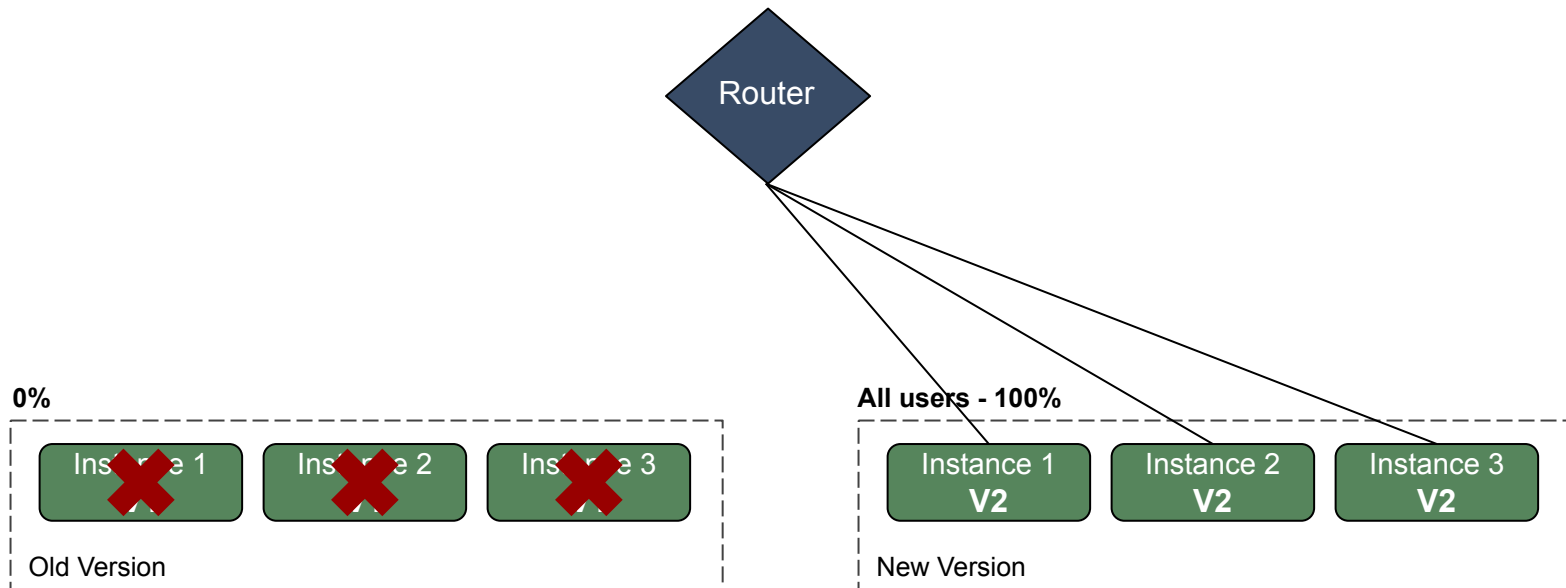
Canary Release



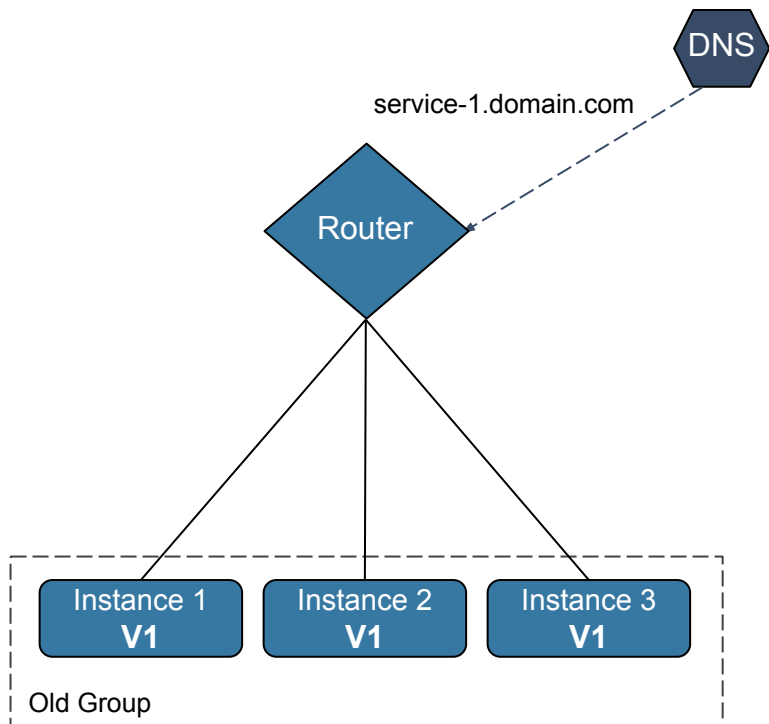
Canary Release



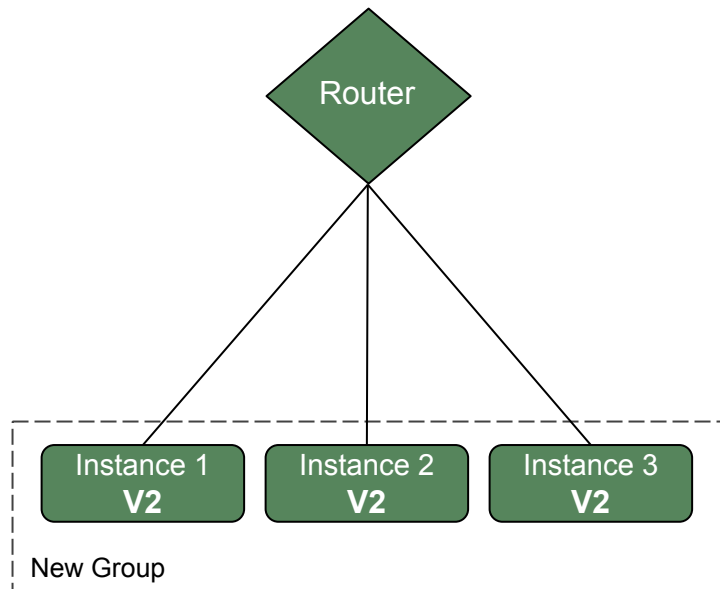
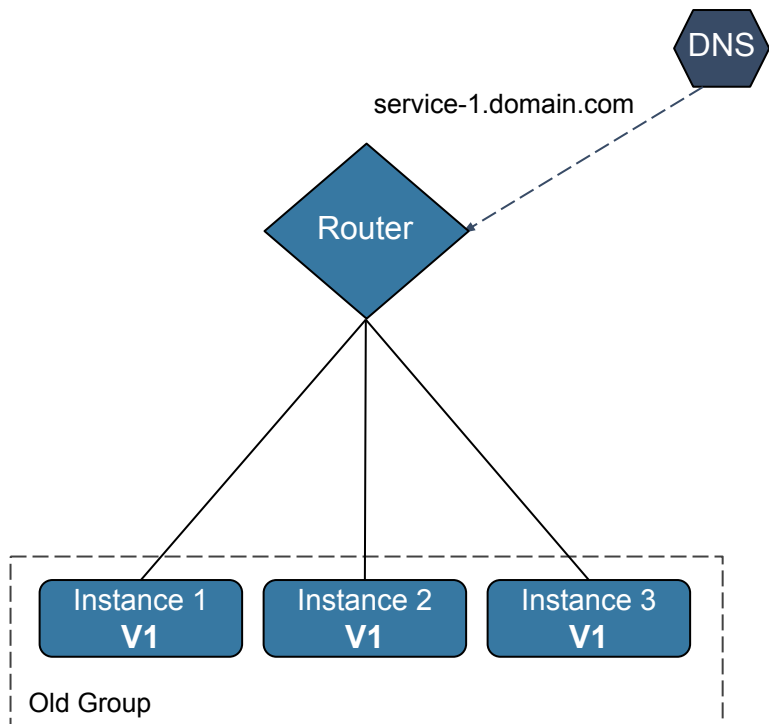
Canary Release



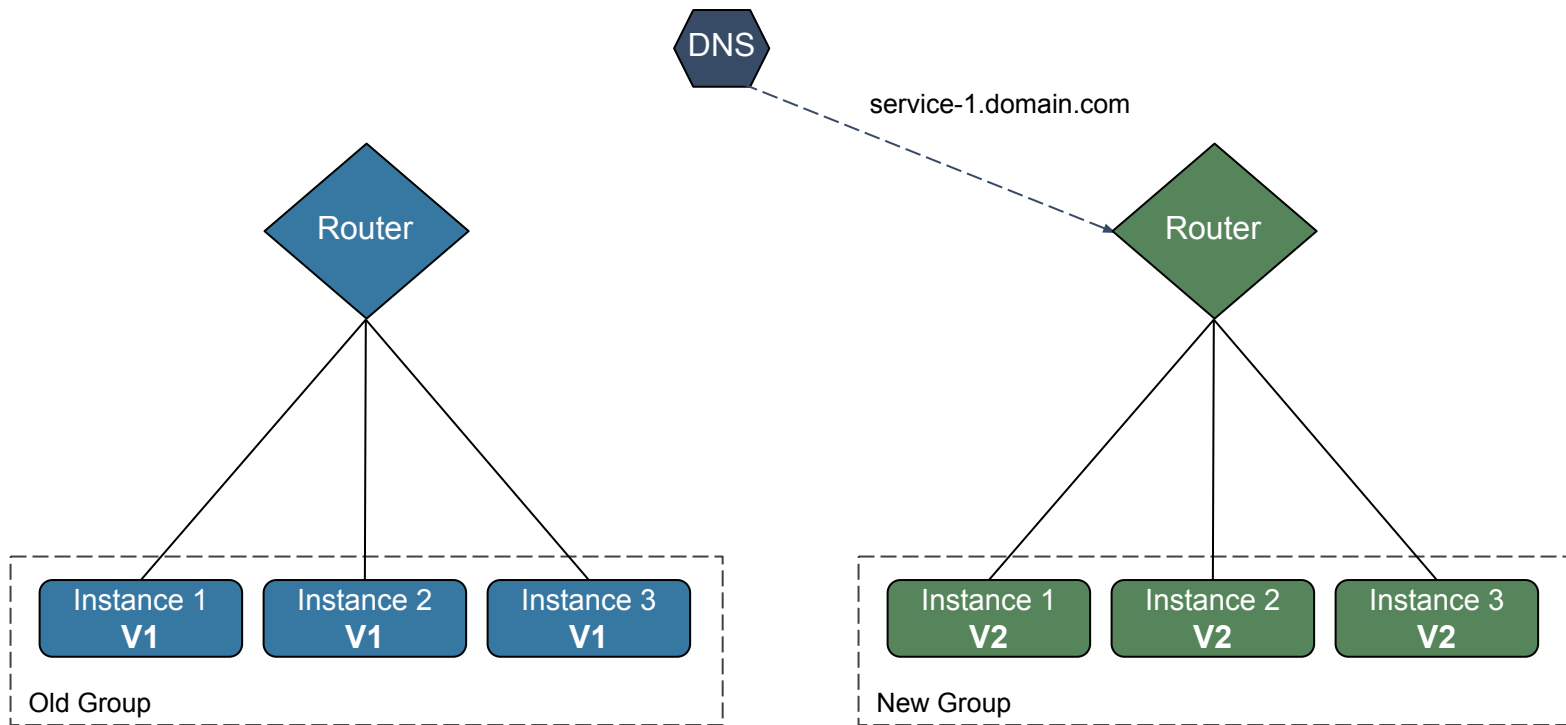
Blue-Green



Blue-Green



Blue-Green

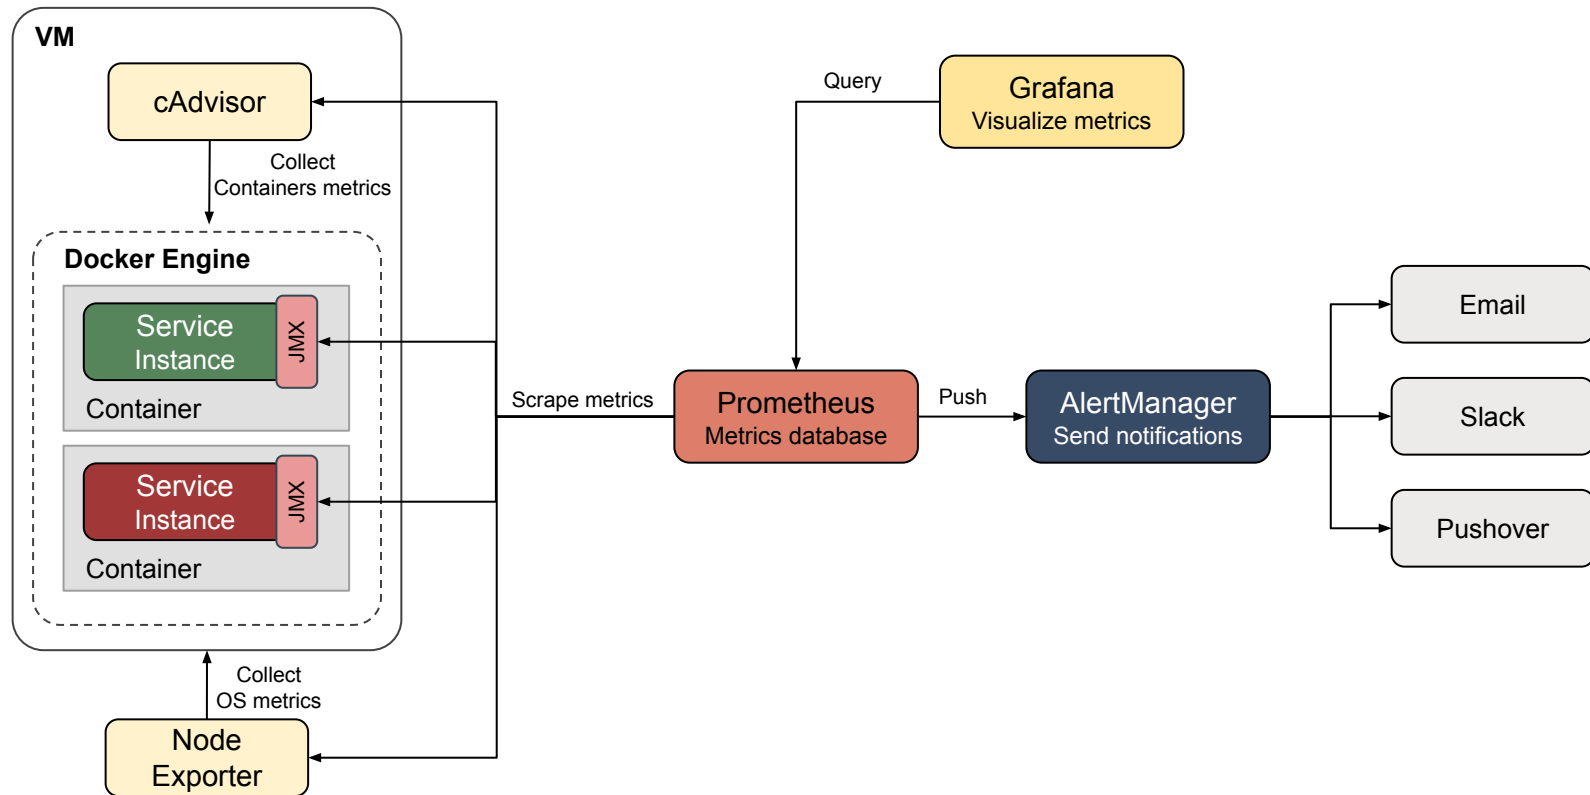


Monitoring Microservices

What to monitor?

- **Hosts**
CPU, Memory, I/O, Network, File system
- **Containers**
CPU, Memory, I/O, Restarts, Bottleneck
- **Applications**
Latency, Traffic, Errors

Metrics & Alerts



Logging & Tracing



Forward the foundation!



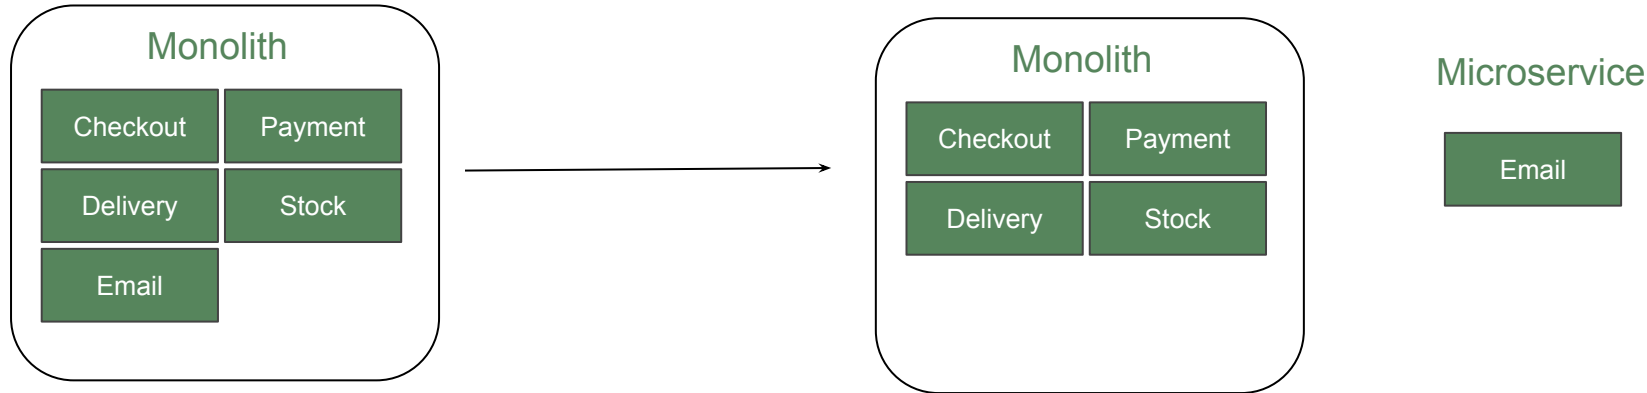
CLOUD NATIVE
COMPUTING FOUNDATION

- Organization that brings together the world's top developers, end users and vendors.
- Hosts critical components of the global technology infrastructure
- Part of the non-profit Linux Foundation

Splitting a Monolith

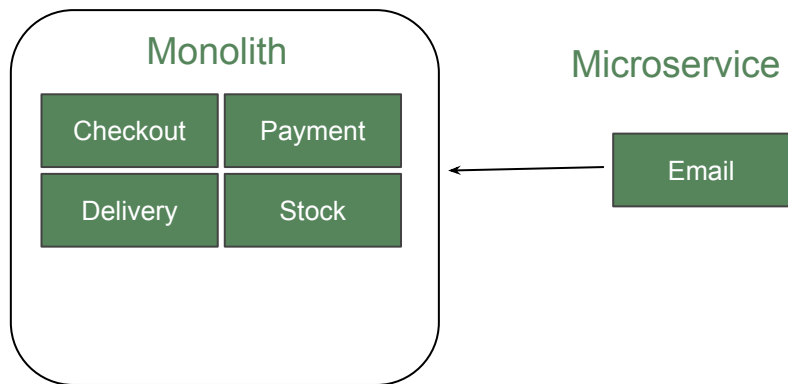
In Microservices

Start with an easy one

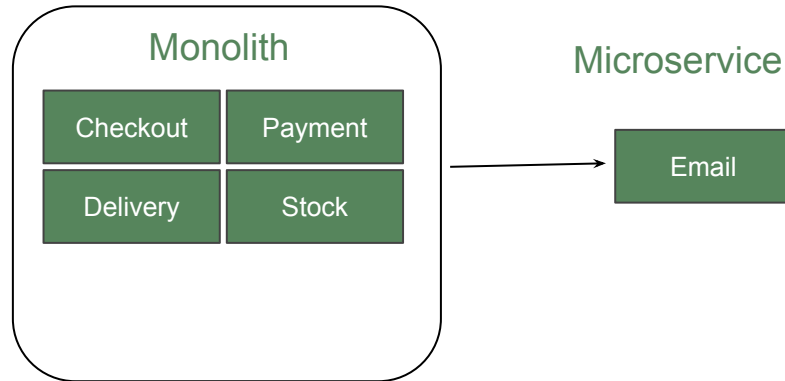


New service should not depend on the monolith

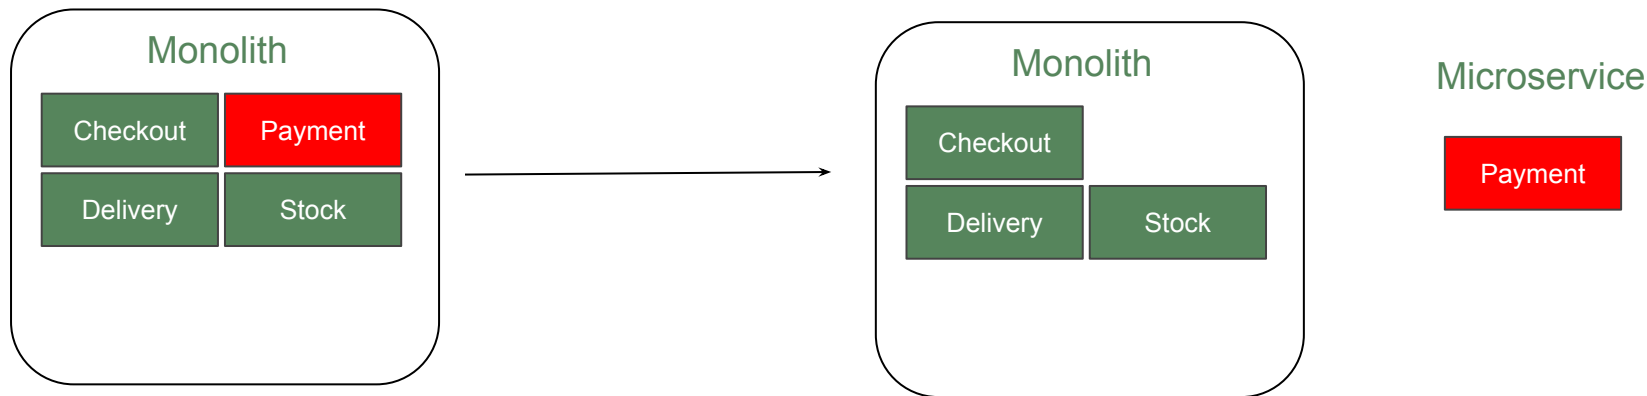
Not okay!



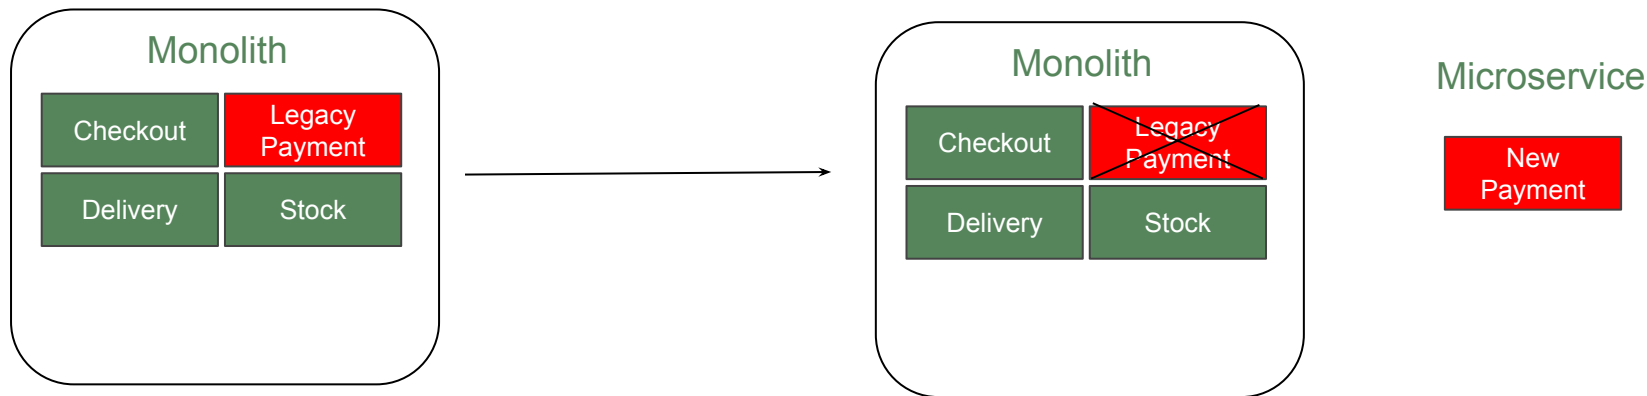
Okay!



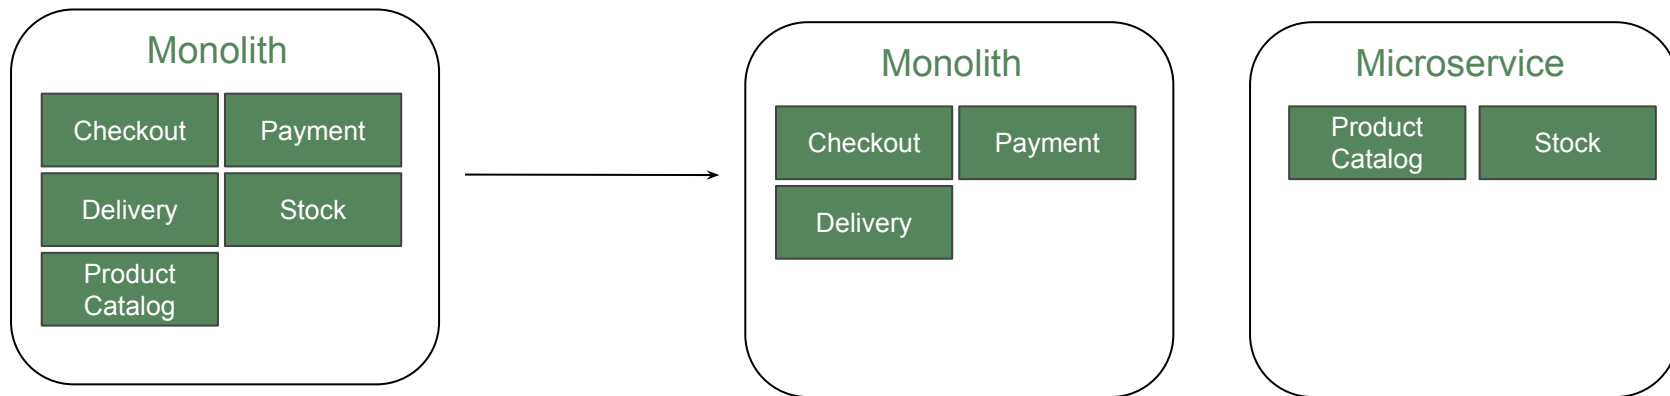
Extract what changes often



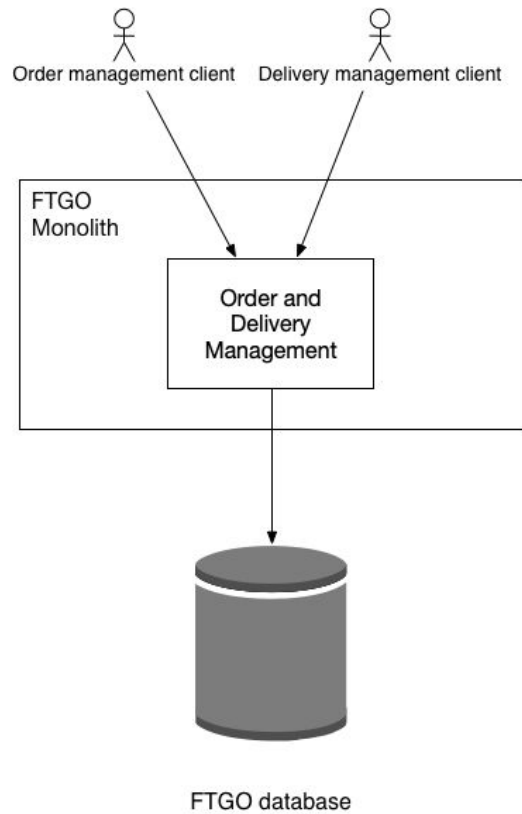
Most of times it is better to rewrite and retire the old code



Better to go from big to small

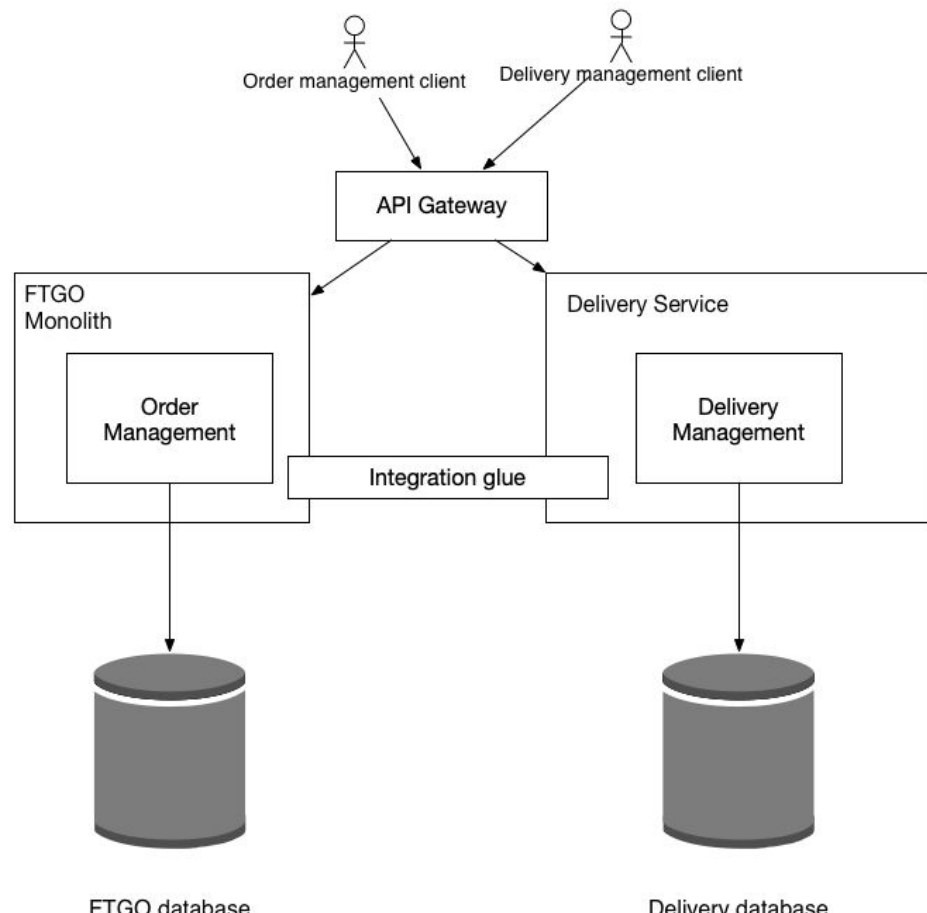


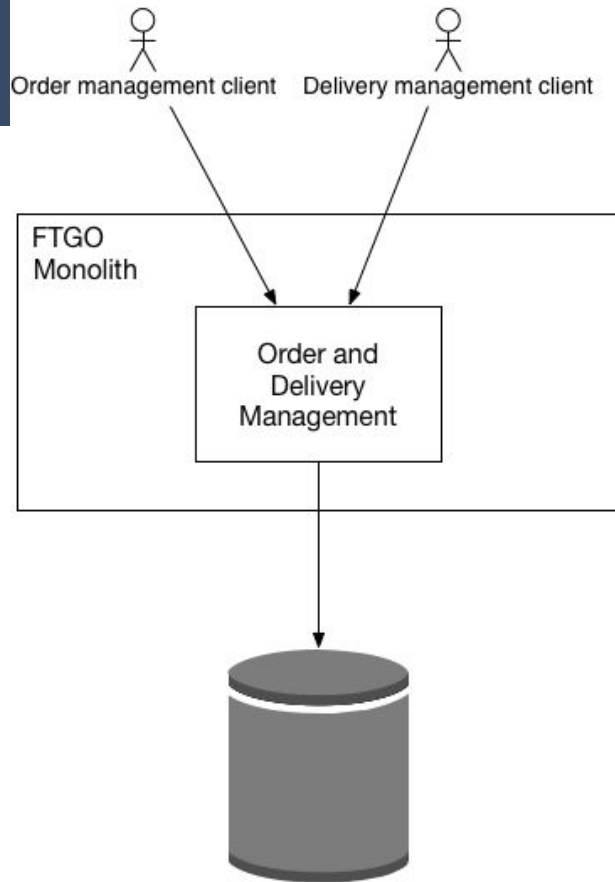
AS-IS



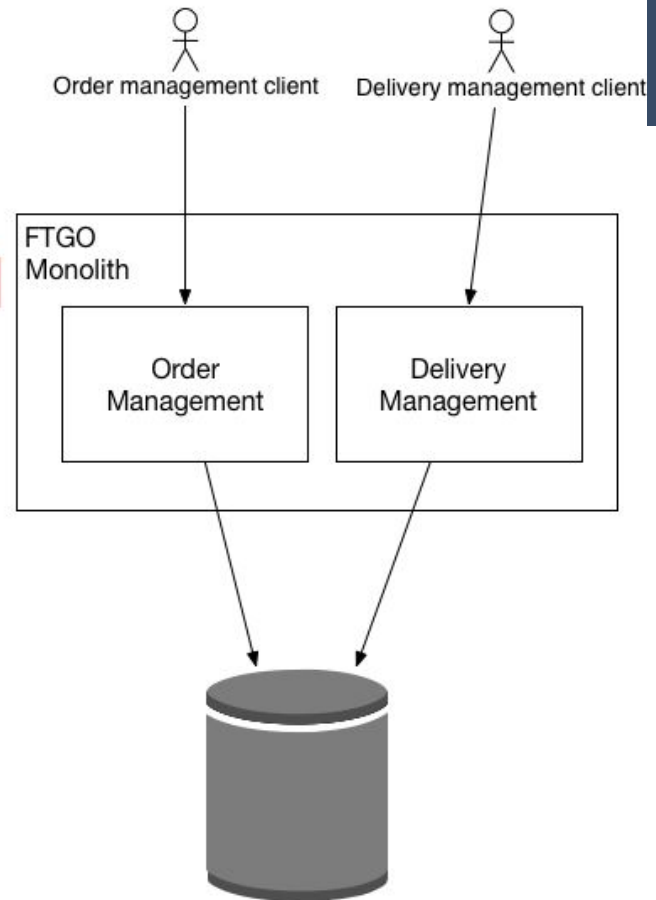
Source: <https://microservices.io/refactoring/>

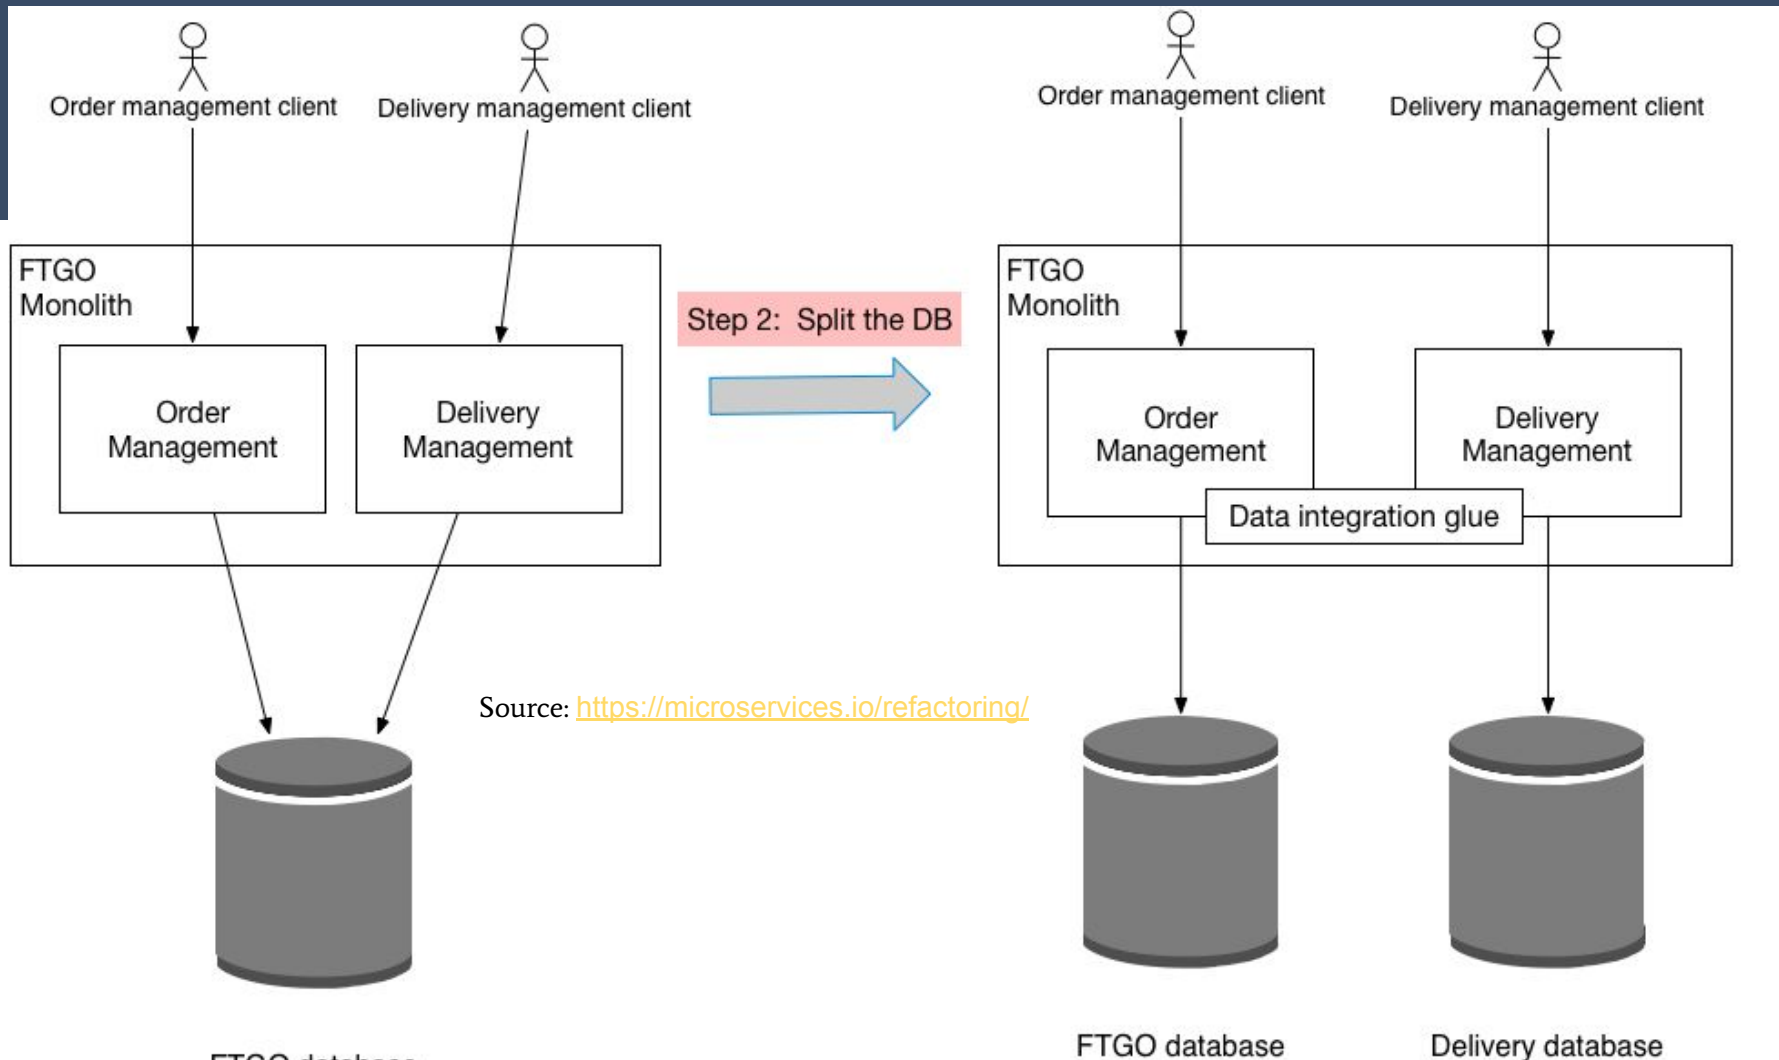
TO-BE

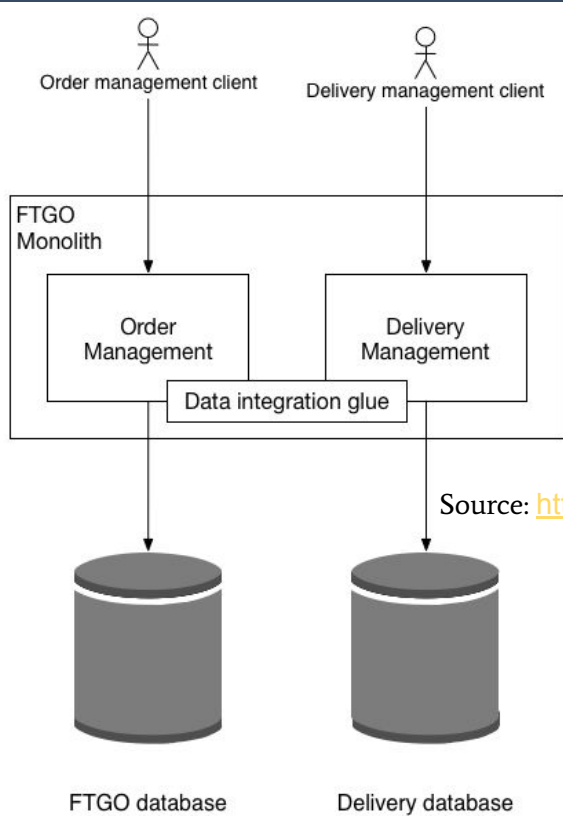




Step 1: Split the code



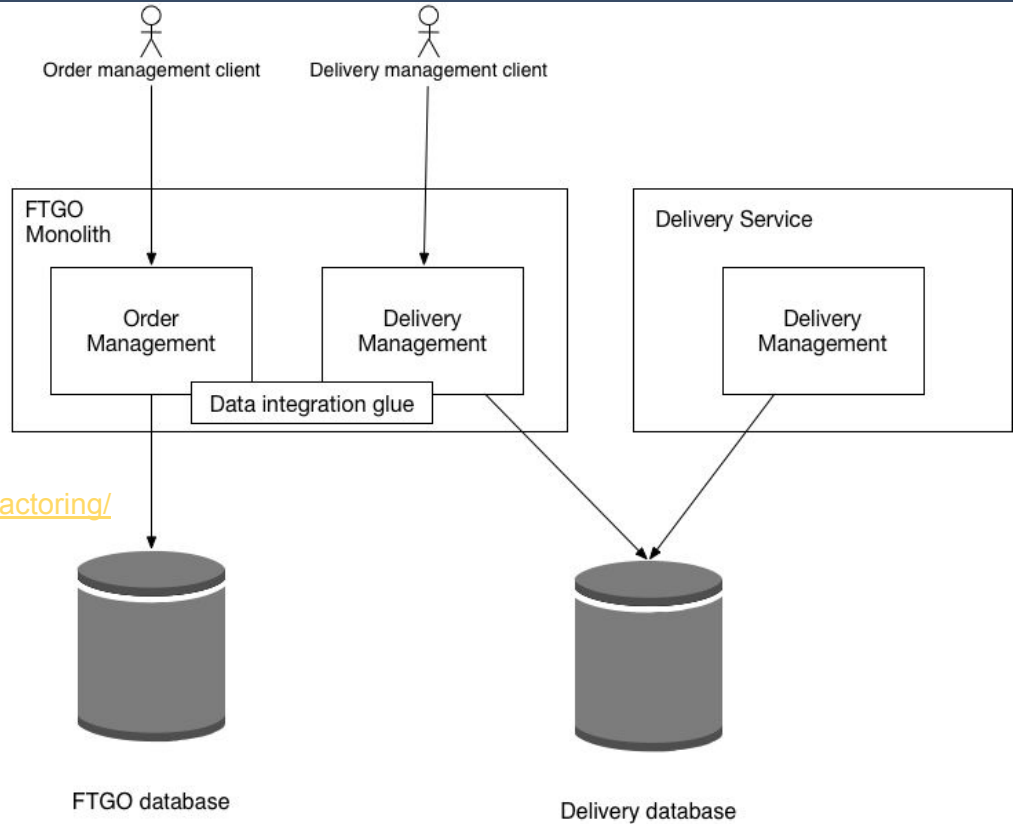


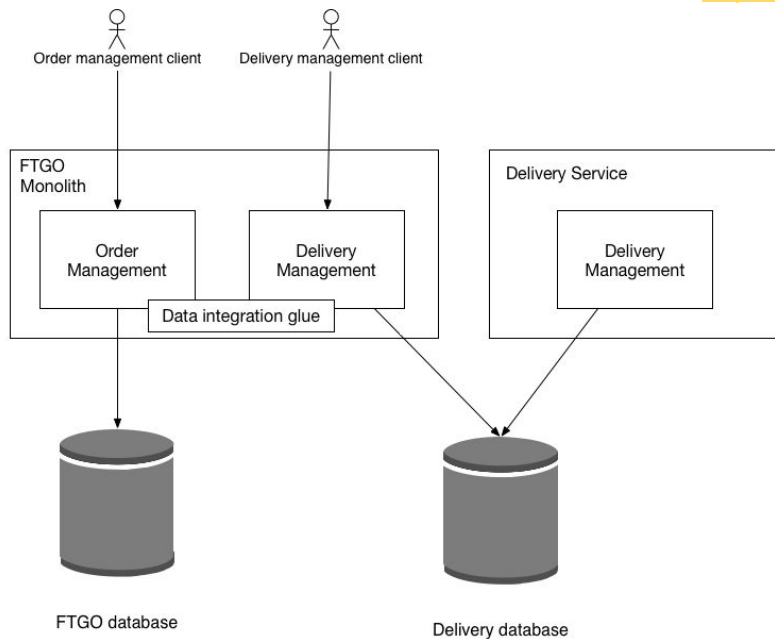


Step 4: Define a service



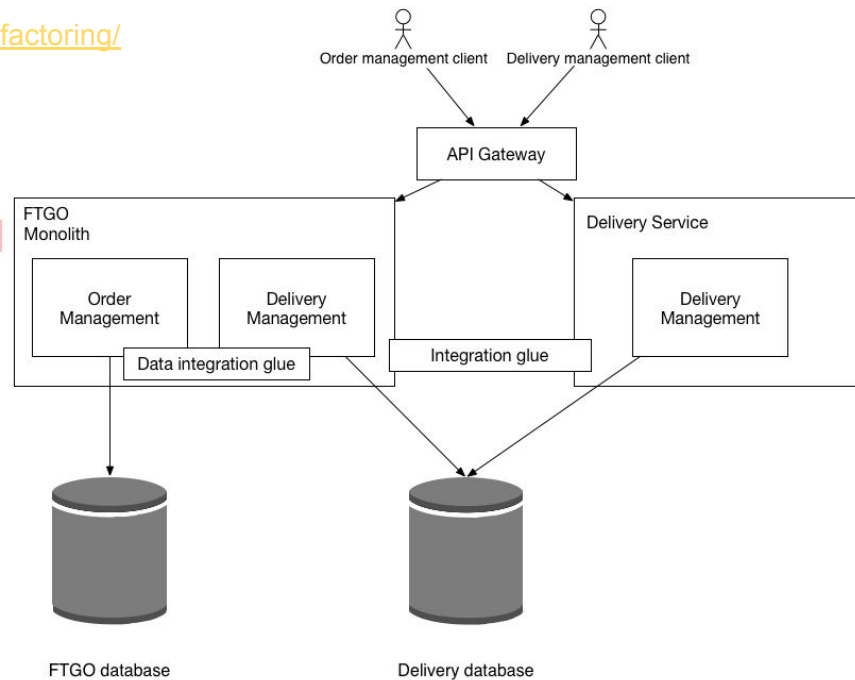
Source: <https://microservices.io/refactoring/>

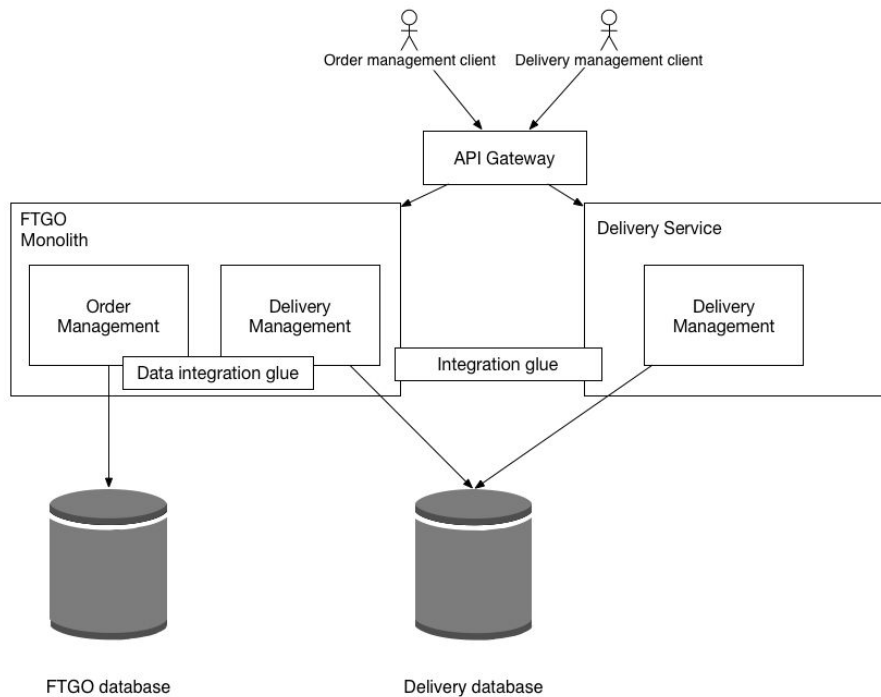




Source: <https://microservices.io/refactoring/>

Step 4: Use Delivery Service





Step 5: Delete old code

